

Deep-Learning:

general principles

+ Convolutional Neural Networks

Pr. Fabien MOUTARDE
Center for Robotics
MINES ParisTech
PSL Université Paris

Fabien.Moutarde@mines-paristech.fr
<http://people.mines-paristech.fr/fabien.moutarde>

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 1

Acknowledgements

During preparation of these slides, I got inspiration and borrowed some slide content from several sources, in particular:

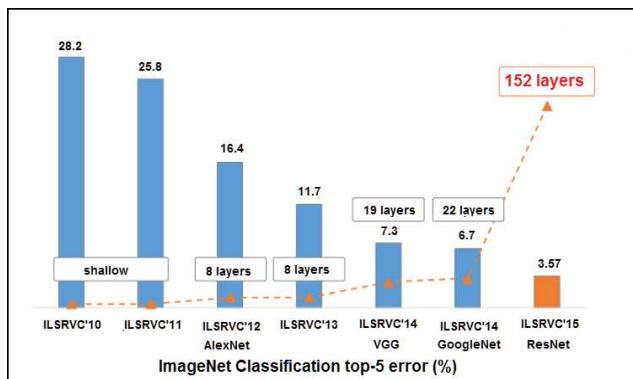
- Yann LeCun + MA Ranzato: slides on « *Deep Learning* » from the corresponding course at NYU
<http://cclv.cs.nyu.edu/doku.php?id=deeplearning:slides:start>
- Hinton+Bengio+LeCun: slides of the *NIPS'2015 tutorial on Deep Learning*
<http://www.iro.umontreal.ca/~bengioy/talks/DL-Tutorial-NIPS2015.pdf>
- Fei-Fei Li + A.Karpathy + J.Johnson: Stanford course lecture slides on « *Convolutional Neural Networks* »
http://cs231n.stanford.edu/slides/winter1516_lecture7.pdf

- **Introduction to Deep Learning**
- **Convolutional Neural Networks (CNN or ConvNets)**
 - Intro + Short reminder on Neural Nets
 - Convolution layers & Pooling layers + global architecture
 - Training algorithm + Dropout Regularization
- **Useful pre-trained convNets**
- **Coding frameworks**
- **Transfer Learning**
- **Object localization and Semantic segmentation**
- **Deep-Learning on 1D signal and 3D data**
- **Recent other image-based applications**

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 3

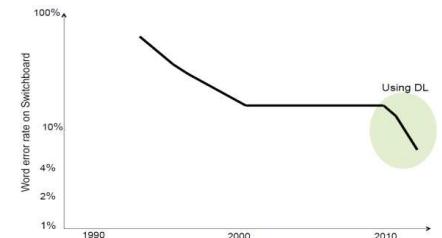
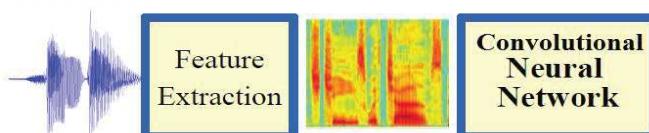
Deep-Learning recent breakthroughs

Very significant improvement over State-of-the-Art in Pattern Recognition / Image Semantic Analysis:



- won many vision pattern recognition competitions (OCR, TSR, object categorization, facial expression,...)
- deployed in photo-tagging by Facebook, Google, Baidu,...

Similar dramatic progress in Speech recognition + Natural Language Processing (NLP)





Object recognition



Scene analysis

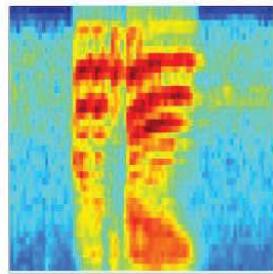


Robotics

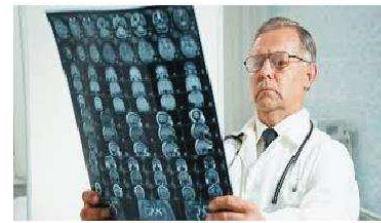
Все счастливые семьи похожи друг на друга, каждая несчастливая семья несчастлива по-своему.

Happy families are all alike.
Every unhappy family is unhappy in its own way.

Language processing



Speech recognition

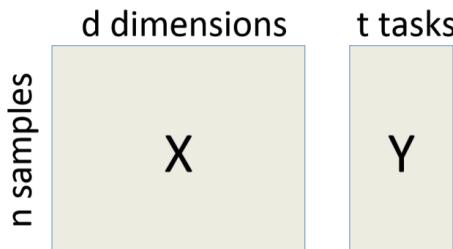


Medical diagnosis & Bio-informatics

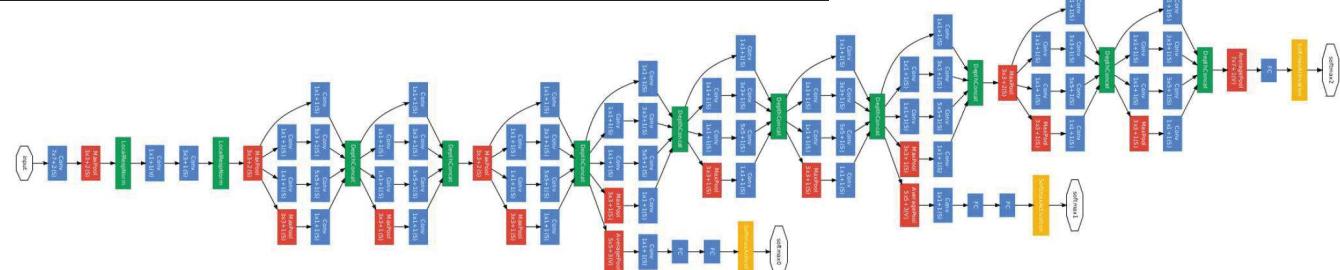
Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 5

Is Deep-Learning « Large-Scale »?

Big and/or « Fat » data



Deep-Learning: Large MODELS



State-of-the-Art Convolutional Neural Networks contain > 100 layers, millions of parameters

Dramatic recent progresses in image classification and visual object categorization not only due to Deep-Learning and convNets:

it was made possible largely thanks to ImageNet dataset, which is a **HUGE** collection of labelled general-purpose images (1000 categories, > 1 million examples)

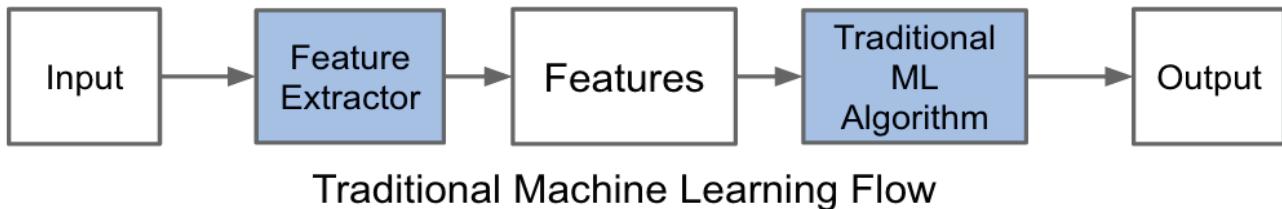
IMAGENET



Most powerful convNets have been trained on this huge dataset!

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 7

Importance of « features » in classical Machine-Learning

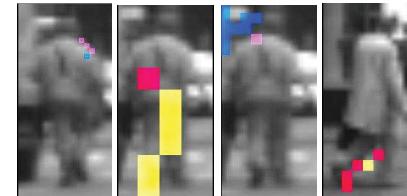


Examples of *hand-crafted* features

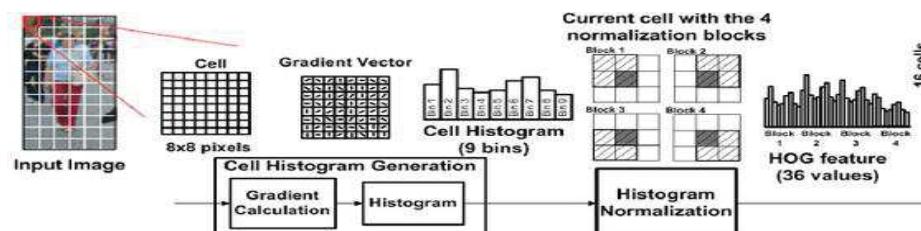
Haar features



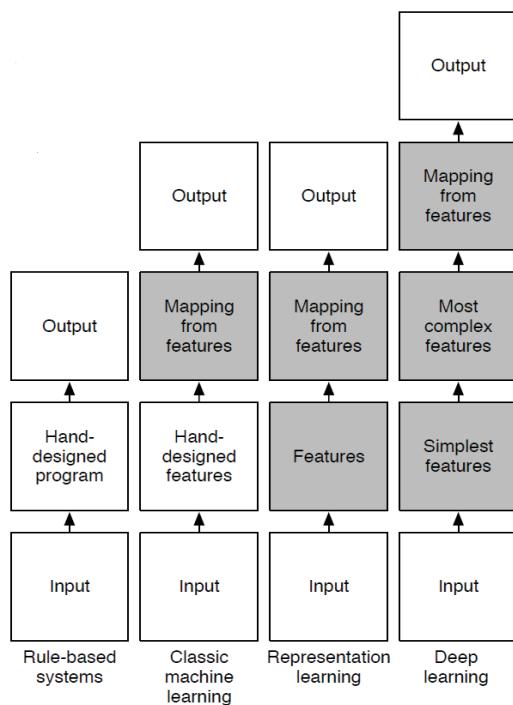
Control-points features



HoG
(Histogram
of Gradients)



Learning a hierarchy of increasingly abstract representations



Increasing level of abstraction
Each stage ~ trainable feature transform

Image recognition

Pixel → edge → texton → motif → part → object

Speech

Sample → spectral band → ... → phoneme → word

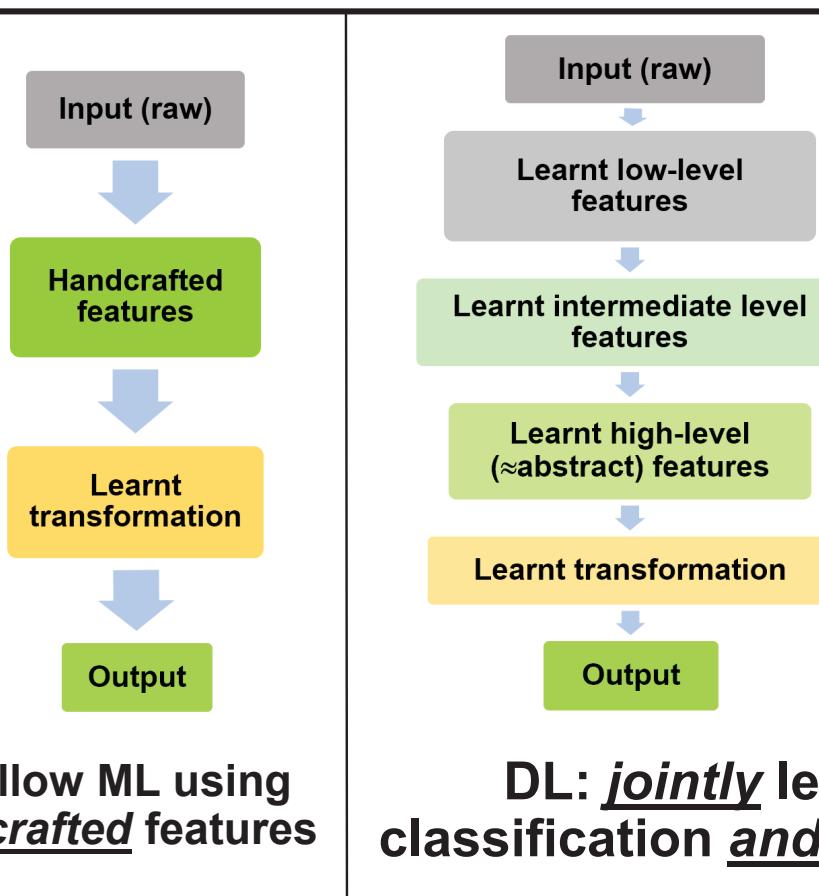
Text

Character → word → word group → clause → sentence → story

[Figure from Goodfellow]

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 9

Deep-Learning vs. shallow Machine-Learning



Why features should be learnt?

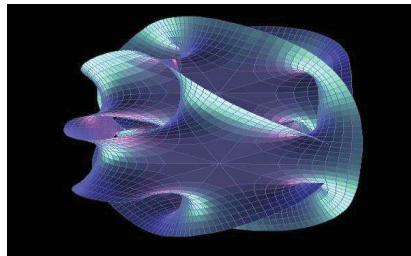
Real data examples for a given task are usually not spreaded everywhere in input space, but rather clustered on a low-dimension « manifold »

Example: Face images of 1000x1000 pixels
→ « raw » examples are vectors in $\mathbb{R}^{1000000}$!!

• BUT:

- position = 3 cartesian coord
- orientation 3 Euler angles
- 50 muscles in face
- Luminosity, color

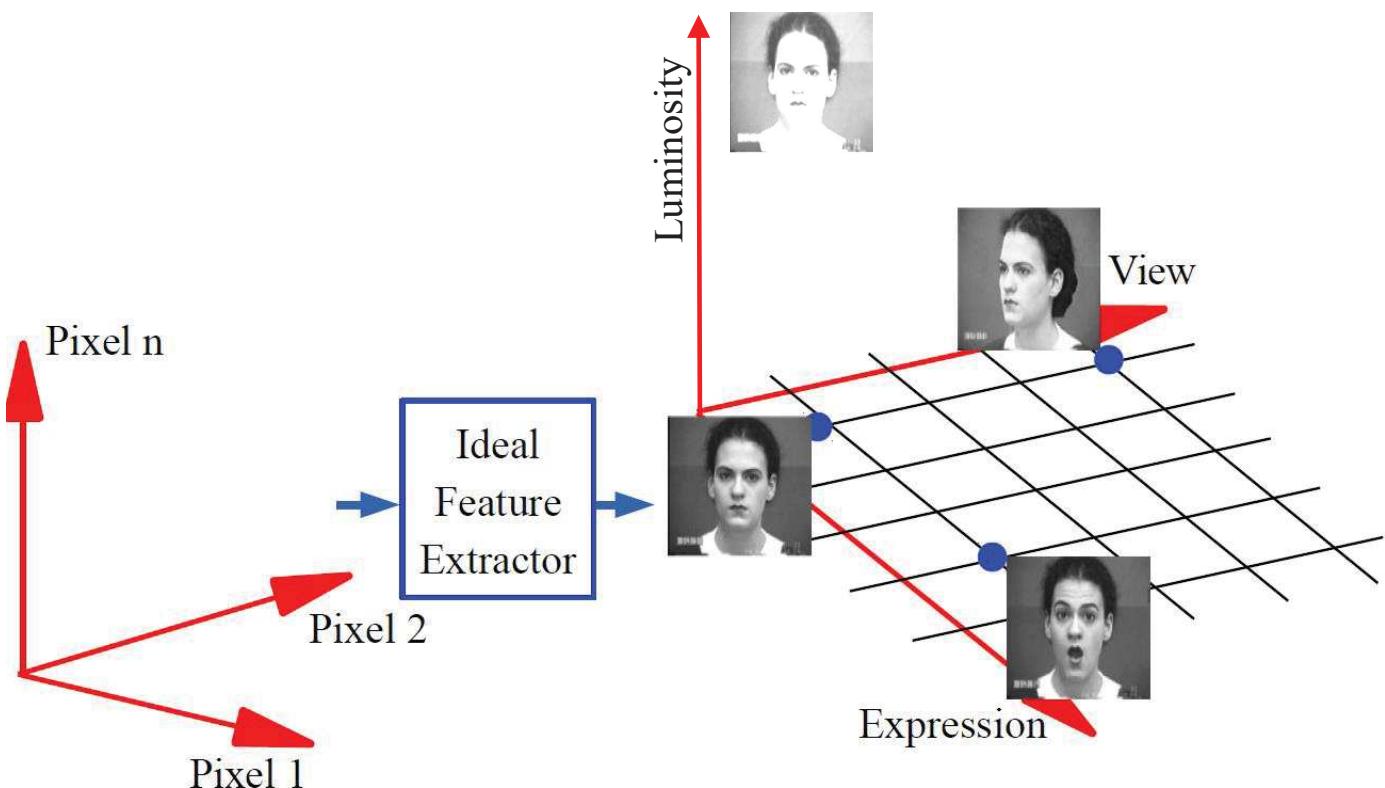
→ Set of all images of ONE person has ≤ 69 dim



→ Examples of face images of 1 person
are all in a LOW-dim manifold
inside a HUGE-dim space

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 11

Good features ~ « mapping » on manifold



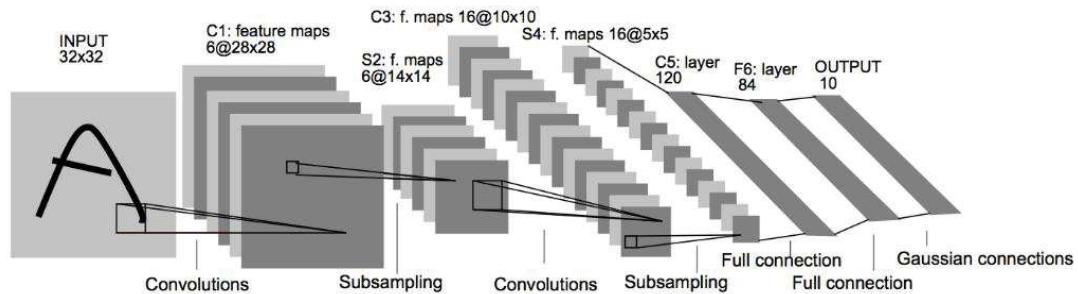
Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 12

- Introduction to Deep Learning
- Convolutional Neural Networks (CNN or ConvNets)
 - Intro + Short reminder on Neural Nets
 - Convolution layers & Pooling layers + global architecture
 - Training algorithm + Dropout Regularization
- Useful pre-trained convNets
- Coding frameworks
- Transfer Learning
- Object localization and Semantic segmentation
- Deep-Learning on 1D signal and 3D data
- Recent other image-based applications

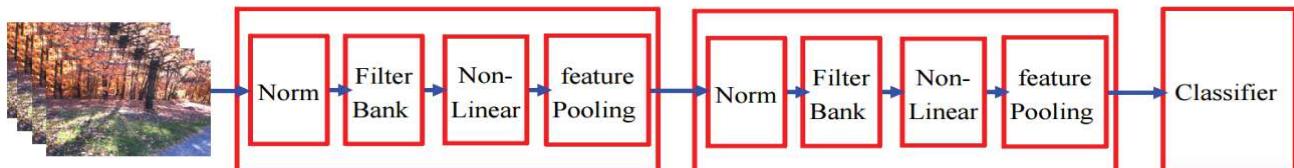
Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 13

Convolutional Neural Networks (CNN, or ConvNet)

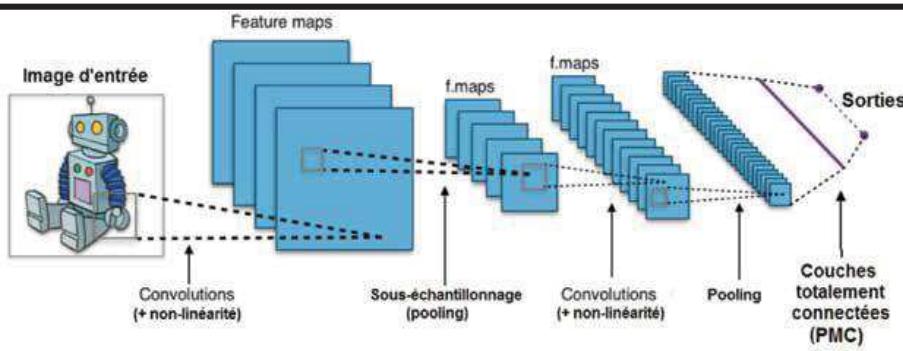
- Proposed in 1998 by Yann LeCun (french prof. @ NYU, now also AI research director of Facebook)



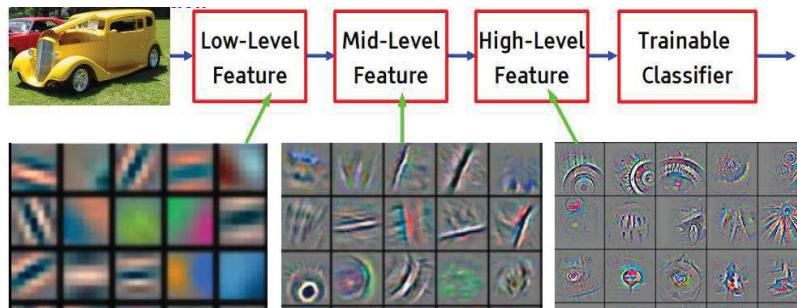
CNN called LeNet by Yann LeCun (1998)



- For inputs with correlated dims (2D image, 1D signal,...)
- Supervised learning

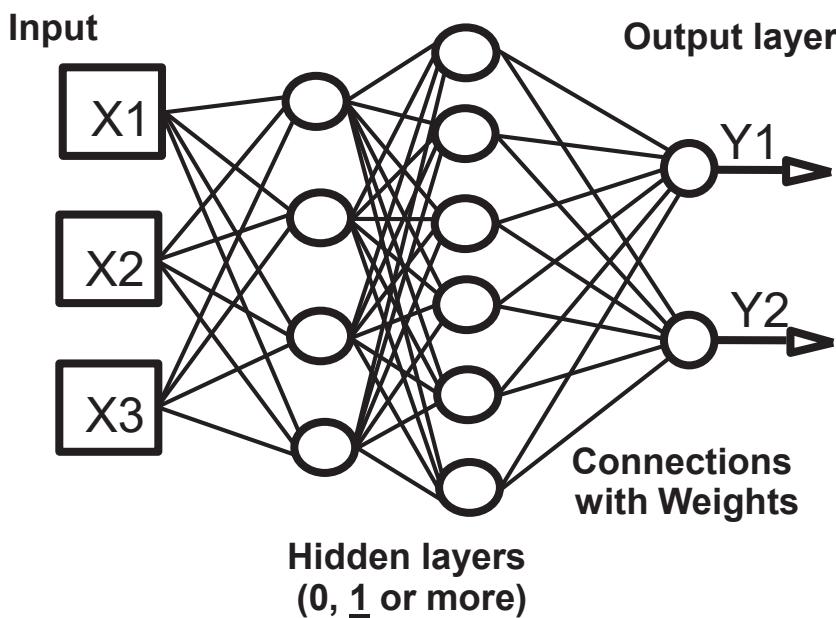


- Wins most vision pattern recognition competitions (OCR, TSR, object categorization, facial expression,...)
- Deployed in photo-tagging by Facebook, Google, Baidu,...
- Also used in real-time video analysis for self-driving cars



Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 15

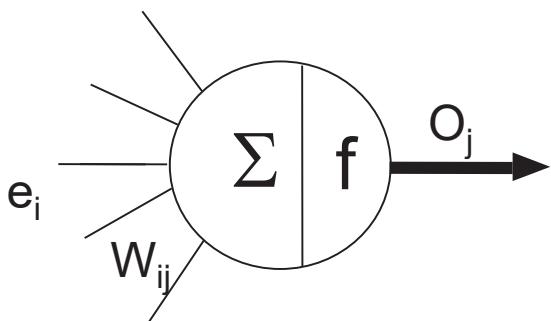
Short reminder on what is a (multi-layer) Neural Network



For “Multi-Layer Perceptron” (MLP),
neurons type generally “summarizing with sigmoid activation”

Reminder on artificial “neurons”

PRINCIPLE

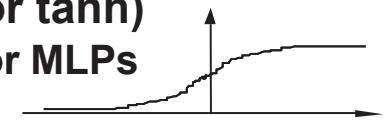


$$O_j = f \left(W_{0j} + \sum_{i=1}^{n_j} W_{ij} e_i \right)$$

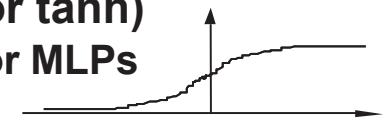
W_{0j} = "bias"

ACTIVATION FUNCTIONS

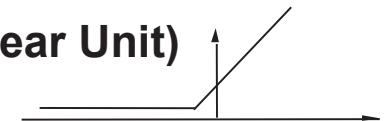
- Threshold (Heaviside or sign)
→ binary neurons



- Sigmoid (logistic or tanh)
→ most common for MLPs



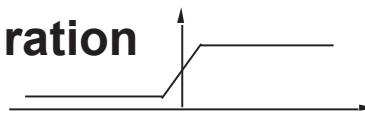
- Identity → linear neurons



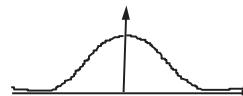
- ReLU (Rectified Linear Unit)



- Saturation

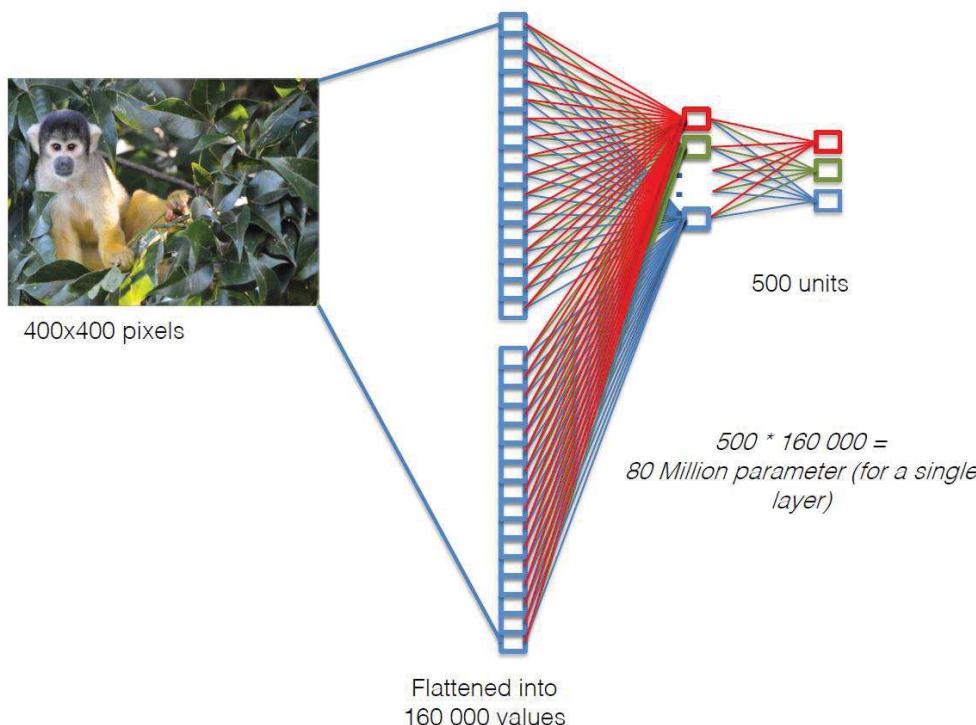


- Gaussian



Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 17

Why MLP directly on pixels is generally a BAD idea?



Huge # of parameters, NO invariance at all

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 18

For image “semantic” classification,
shift-invariance of features is useful



And ANY shift-invariant & linear system can always
be expressed as a CONVOLUTION:

$$y[n] = \sum x[m] h[n-m]$$

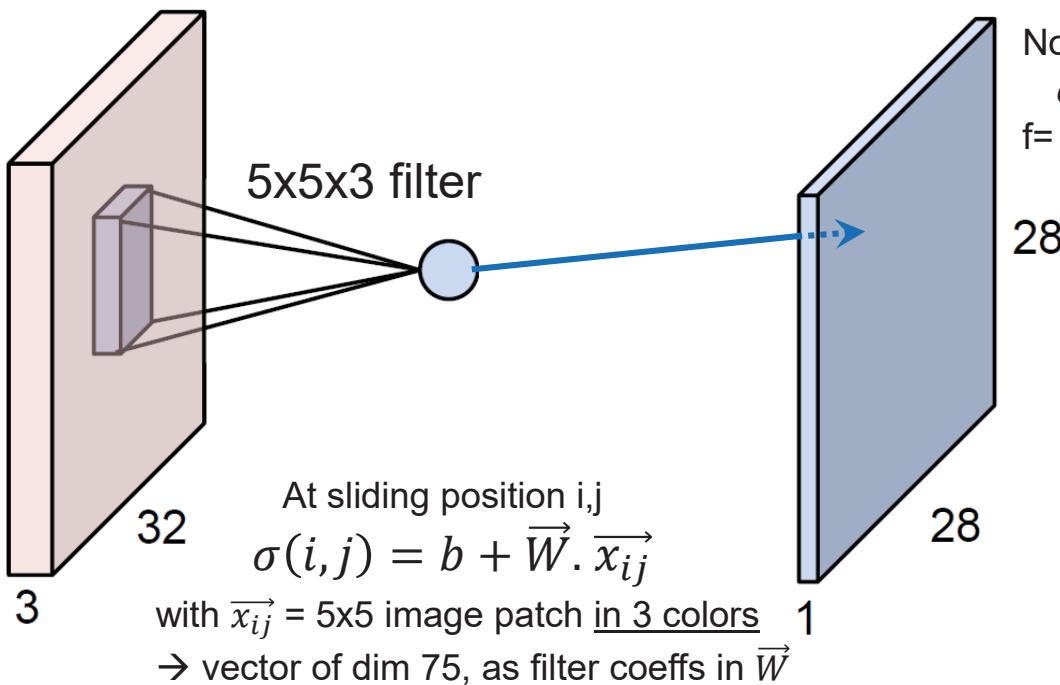
(where $h[n]$ is the impulse response).

Outline

- Introduction to Deep Learning
- Convolutional Neural Networks (CNN or ConvNets)
 - Intro + Short reminder on Neural Nets
 - Convolution layers & Pooling layers + global architecture
 - Training algorithm + Dropout Regularization
- Useful pre-trained convNets
- Coding frameworks
- Transfer Learning
- Object localization and Semantic segmentation
- Deep-Learning on 1D signal and 3D data
- Recent other image-based applications

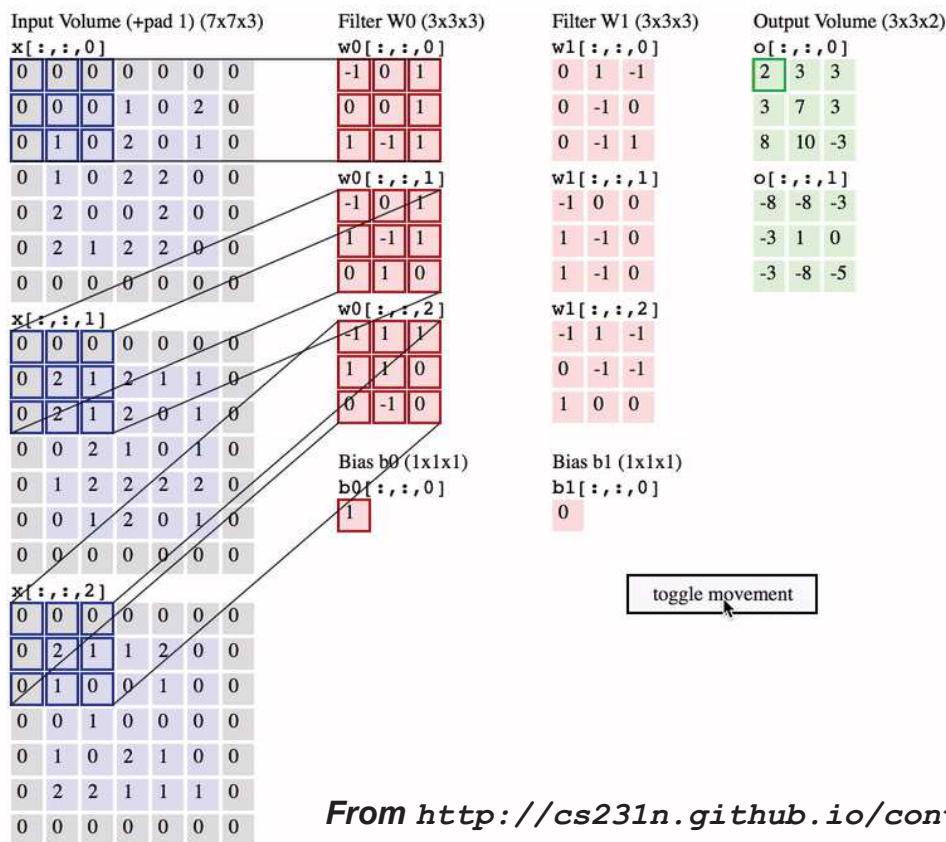
32x32x3 image

activation map

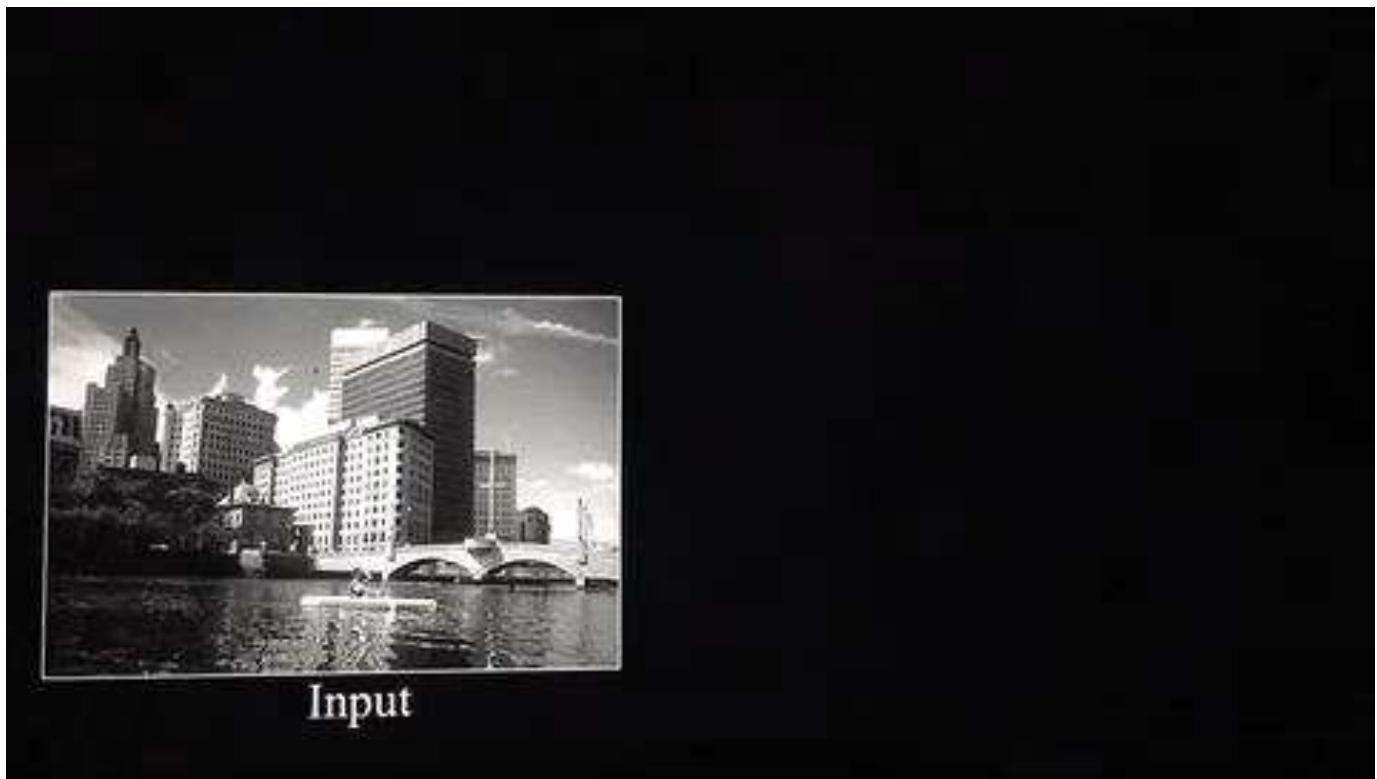


Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 21

Convolution in action

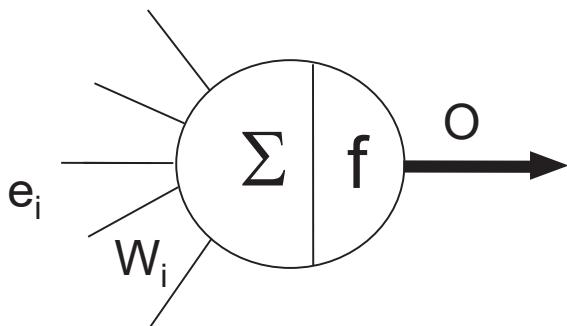


From <http://cs231n.github.io/convolutional-networks/>



Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 23

« Neural » view of convolution filters and layers

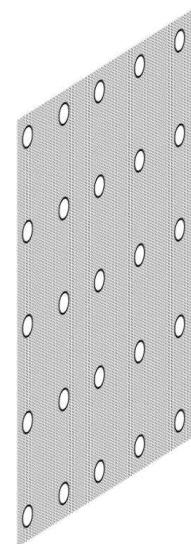


$$O = f\left(W_0 + \sum_{i=1}^n W_i e_i\right)$$

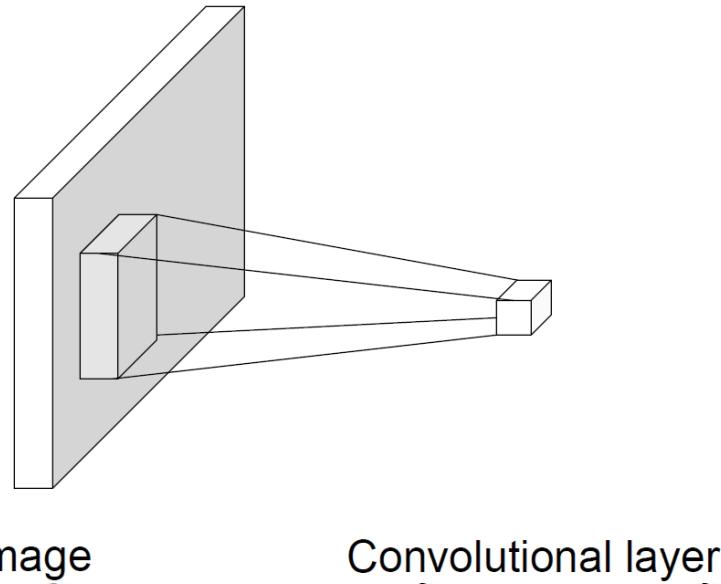
W_0 = "bias"

f = activation function

Each convolution FILTER is one set of neuron parameters

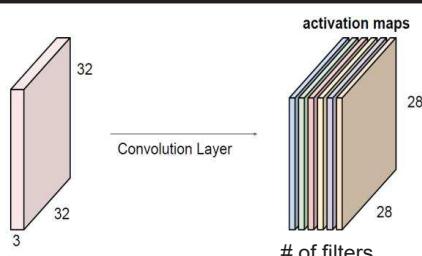


Each convolution LAYER is a set of ~imageSize neurons, but they all have same SHARED weights (perform SAME convolution)



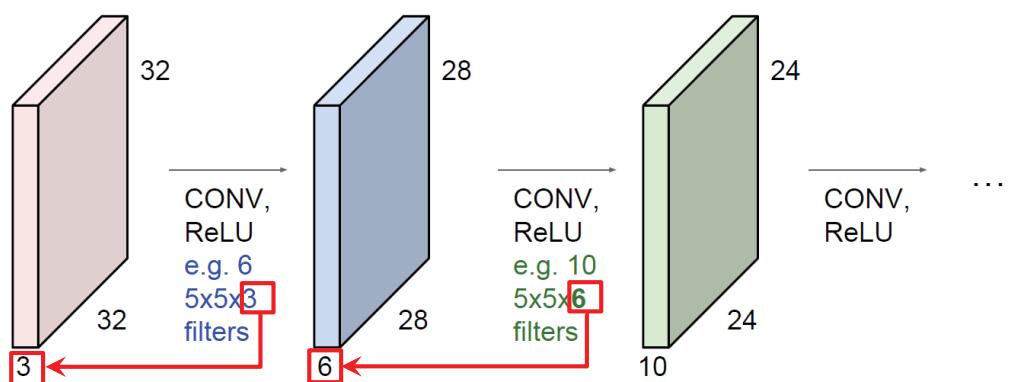
Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 25

Convolutional layers

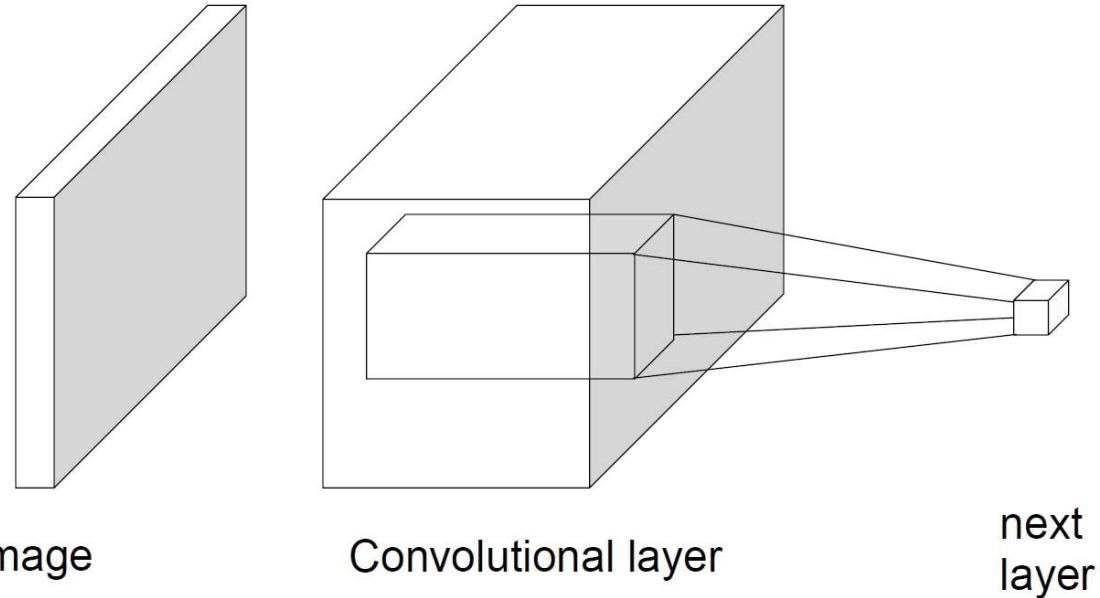


One “activation map” for each convolution filter

A convNet: succession of Convolution+activation Layers



NB: each convolution layer processes FULL DEPTH of previous activation map

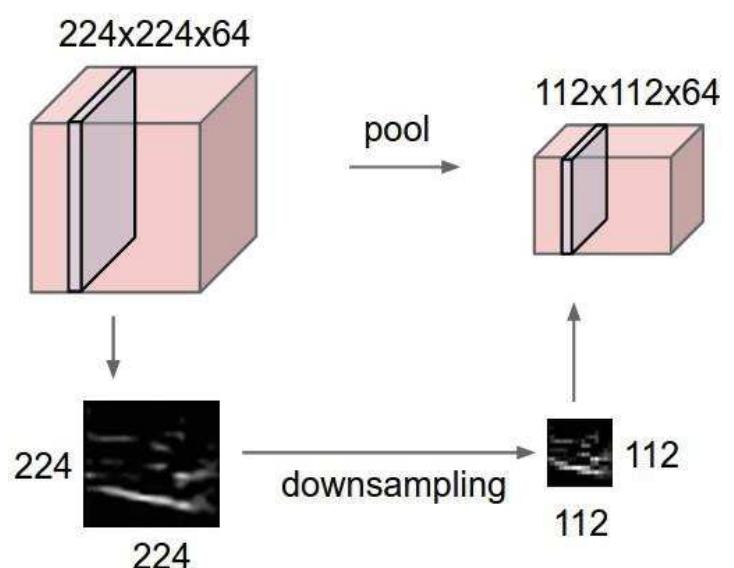


Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 27

Pooling layers

Goal:

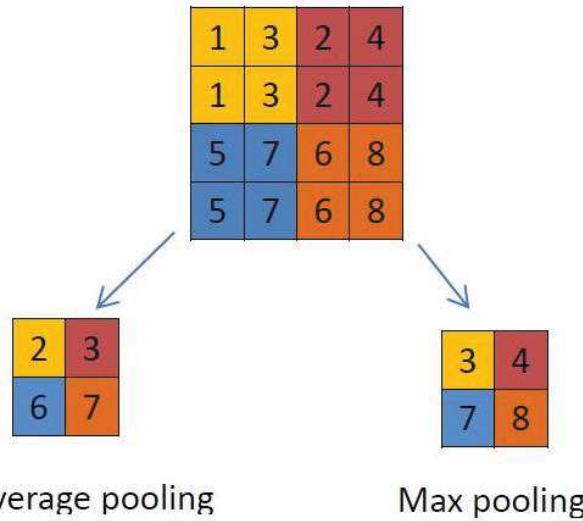
- aggregation over space
- noise reduction,
- small-translation invariance,
- small-scaling invariance



Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 28

Parameters:

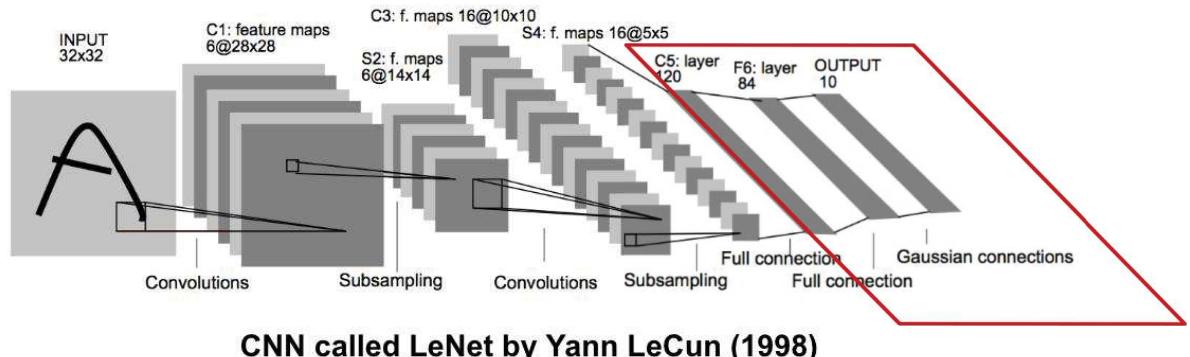
- pooling size (often 2×2)
- pooling stride (usually = pooling_size)
- Pooling operation: max, average, Lp,...



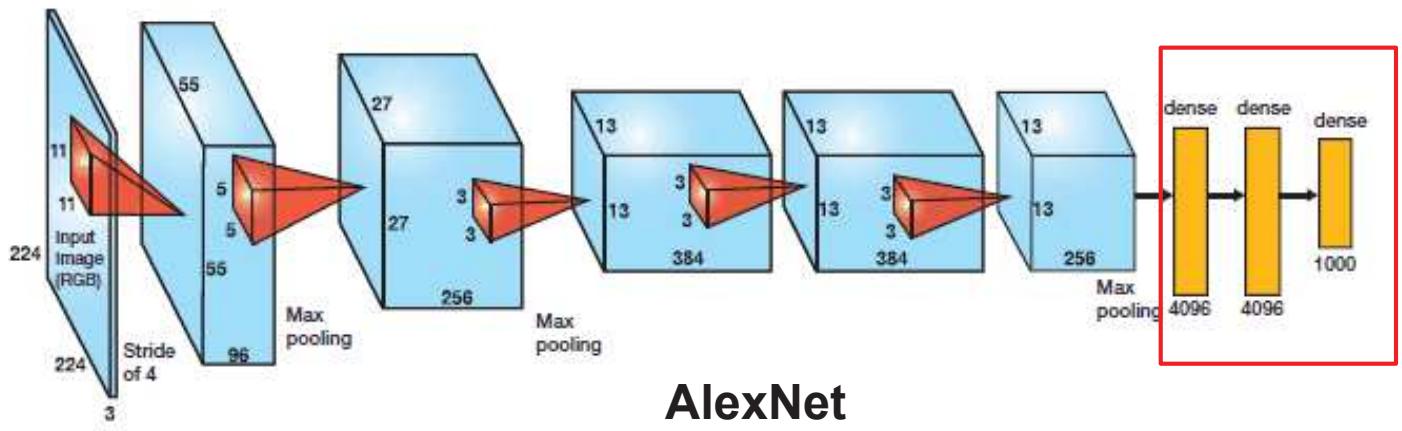
Example: 2×2 pooling, stride 2

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 29

Final classification layer: just a classical MLP

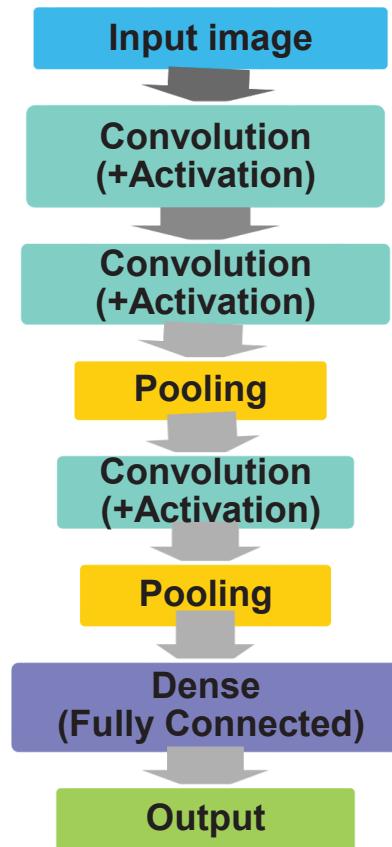


CNN called LeNet by Yann LeCun (1998)



AlexNet

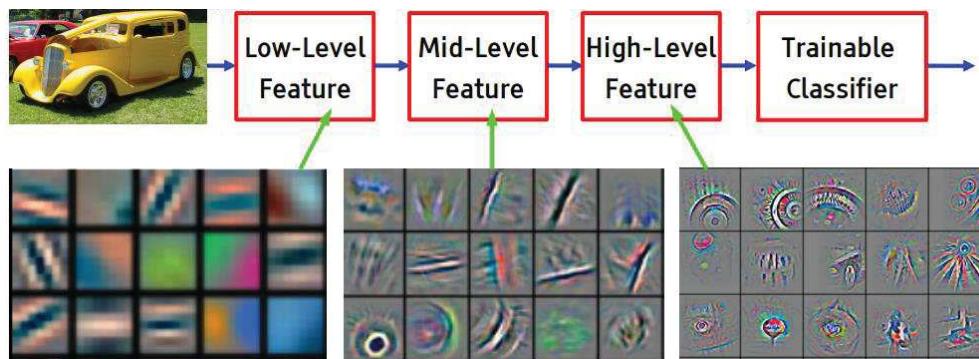
Succession of Convolution (+ optional activation) layers and Pooling layers, which extract the hierarchy of features, followed by dense (fully connected) layer(s) for final classification



Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 31

Typical convolutional filters after training

Architecture with a deep succession of layers processing coarser and coarser “images”
→ Lower layer learns optimized low-level filters
(detection of ~edges in L1, ~corners/arcs in L2)
→ Higher level layers learn more “abstract” filters
(~“texture types” in L3, ~object parts in L4)
→ Last layer output a representation on which it is easy to discriminate between classes



Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 32

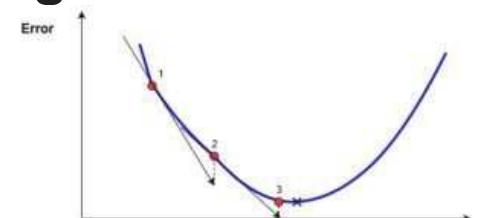
- Introduction to Deep Learning
- Convolutional Neural Networks (CNN or ConvNets)
 - Intro + Short reminder on Neural Nets
 - Convolution layers & Pooling layers + global architecture
 - Training algorithm + Dropout Regularization
- Useful pre-trained convNets
- Coding frameworks
- Transfer Learning
- Object localization and Semantic segmentation
- Deep-Learning on 1D signal and 3D data
- Recent other image-based applications

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 33

ConvNet training

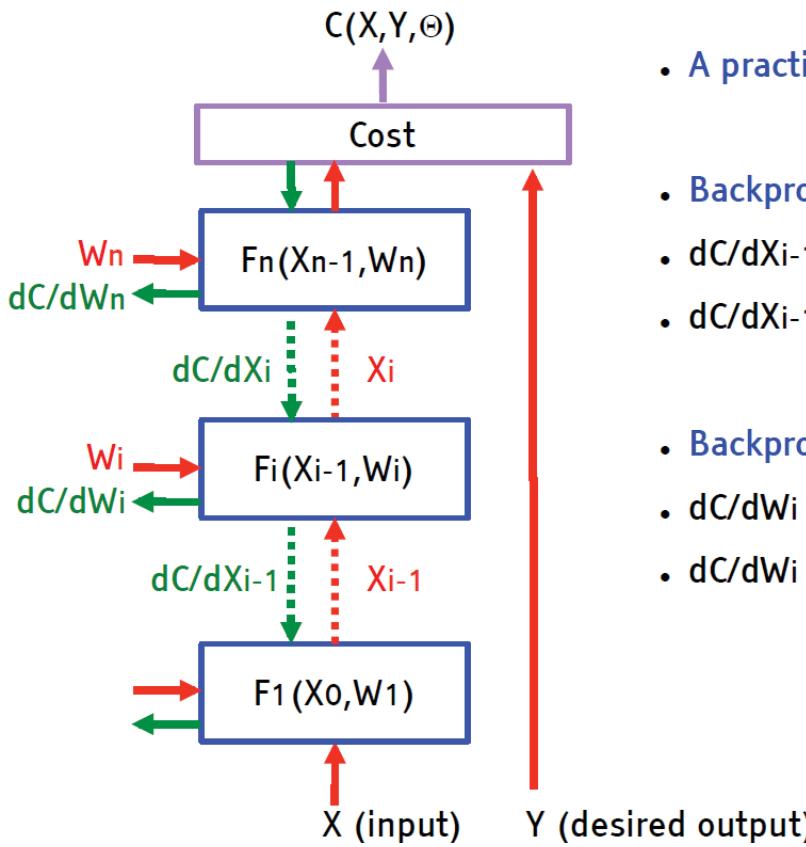
All successive layers of a convNet forms a Deep neural network (with weigh-sharing inside each conv. Layer, and specific pooling layers).

Training = optimizing values of weights&biases
Method used = gradient descent



➔ Stochastic Gradient Descent (SGD), using back-propagation:

- Input 1 (or a few) random training sample(s)
- Propagate
- Calculate error (loss)
- Back-propagate through all layers from end to input, to compute gradient
- Update convolution filter weights



- A practical Application of Chain Rule

- Backprop for the state gradients:

- $dC/dX_{i-1} = dC/dX_i \cdot dX_i/dX_{i-1}$
- $dC/dX_{i-1} = dC/dX_i \cdot dF_i(X_{i-1}, W_i)/dX_{i-1}$

- Backprop for the weight gradients:

- $dC/dW_i = dC/dX_i \cdot dX_i/dW_i$
- $dC/dW_i = dC/dX_i \cdot dF_i(X_{i-1}, W_i)/dW_i$

Recall of back-prop principle

Smart method for efficient computing of gradient (w.r.t. weights) of a Neural Network cost function, based on chain rule for derivation.

Cost function is $Q(t) = \sum_m \text{loss}(Y_m, D_m)$, where m runs over training set examples

Usually, $\text{loss}(Y_m, D_m) = ||Y_m - D_m||^2$ [quadratic error]

Total gradient:

$$W(t+1) = W(t) - \lambda(t) \text{grad}_W(Q(t)) + \mu(t)(W(t) - W(t-1))$$

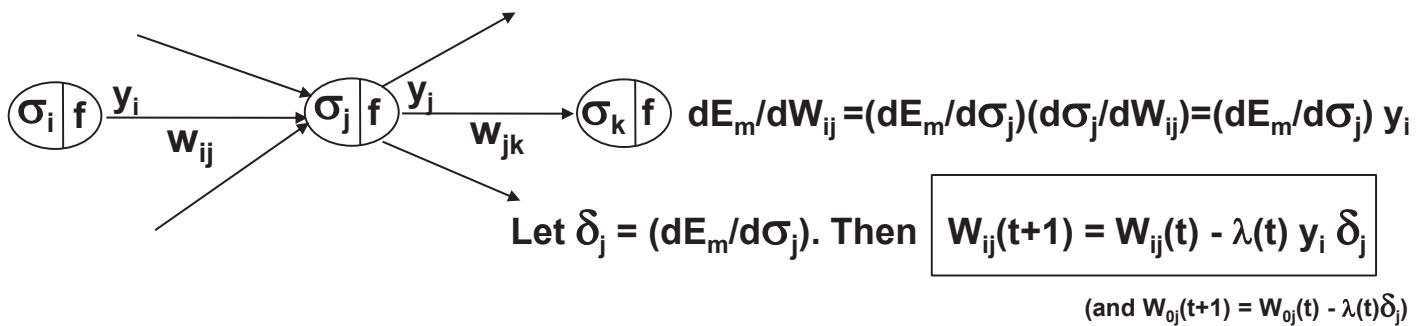
Stochastic gradient:

$$W(t+1) = W(t) - \lambda(t) \text{grad}_W(Q_m(t)) + \mu(t)(W(t) - W(t-1))$$

where $Q_m = \text{loss}(Y_m, D_m)$, is error computed on only ONE example randomly drawn from training set at every iteration and

$\lambda(t)$ = learning rate (fixed, decreasing or adaptive), $\mu(t)$ = momentum

Now, how to compute dQ_m/dW_{ij} ?



If neuron j is output, $\delta_j = (dE_m/d\sigma_j) = (dE_m/dy_j)(dy_j/d\sigma_j)$ with $E_m = ||Y_m - D_m||^2$

so $\delta_j = 2(y_j - D_j)f'(\sigma_j)$ if neuron j is an output

Otherwise, $\delta_j = (dE_m/d\sigma_j) = \sum_k (dE_m/d\sigma_k)(d\sigma_k/d\sigma_j) = \sum_k \delta_k (d\sigma_k/d\sigma_j) = \sum_k \delta_k w_{jk} (dy_j/d\sigma_j)$

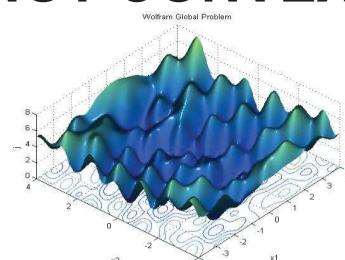
so $\delta_j = (\sum_k w_{jk} \delta_k) f'(\sigma_j)$ if neuron j is "hidden"

→ all the δ_j can be computed successively from last layer to upstream layers by "error backpropagation" from output

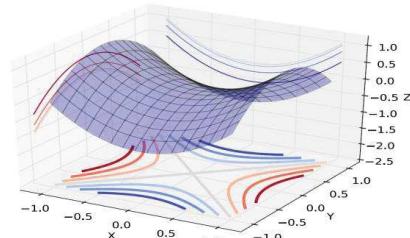
Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 37

Why gradient descent works despite non-convex?

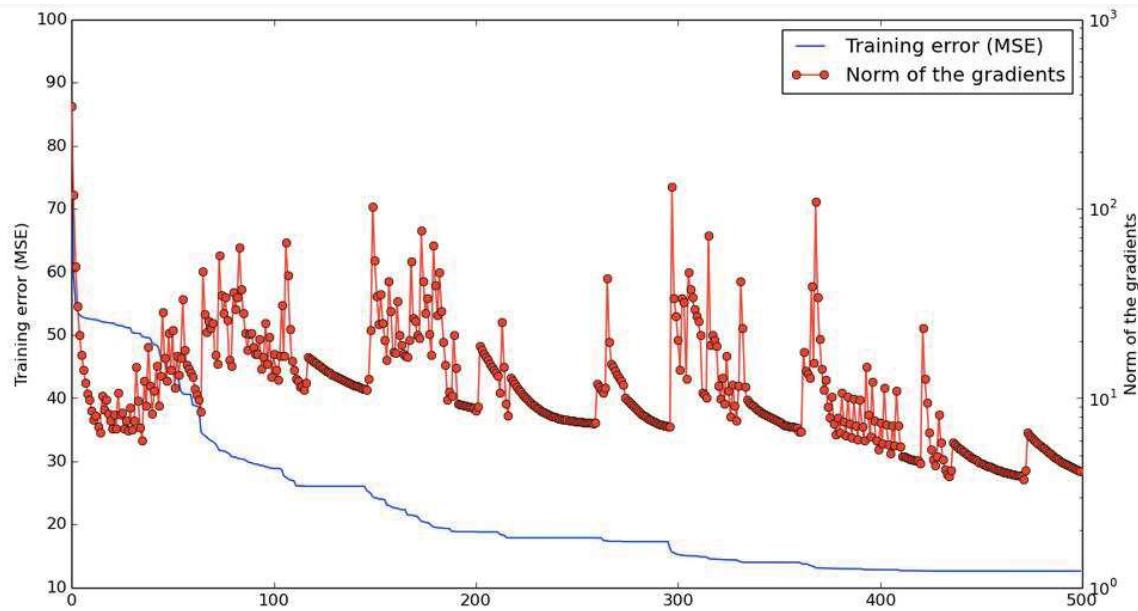
Error surface for neural net are NOT CONVEX !



- Local minima dominate in low-Dim...
- ...but recent work has shown that saddle points dominate in high-Dim



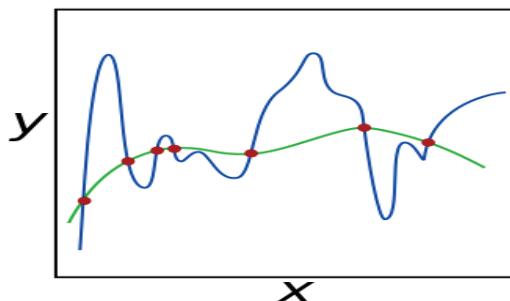
- Furthermore, most local minima are close to the global minimum



- Oscillating between two behaviors:
 - Slowly approaching a saddle point
 - Escaping it

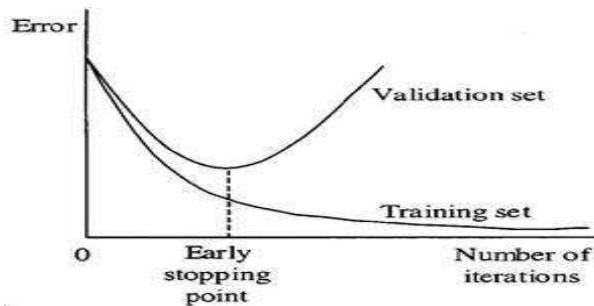
Some ConvNet training « tricks »

- Importance of input normalization
(zero mean, unit variance)
- Importance of weights initialization
random but **SMALL** and prop. to $1/\sqrt{\text{nbInputs}}$
- Decreasing (or adaptive) learning rate
- Importance of training set size
ConvNets often have a **LARGE** number of free parameters
→ train them with a sufficiently large training-set !
- Avoid overfitting by:
 - Early Stopping of learning iterations
 - Use of L1 or L2 regularization (after some epochs)
 - Use « Dropout » regularization (esp. on large FC layers)



Trying to fit too many free parameters with not enough information can lead to *overfitting*

How to detect overfitting for iterative training?

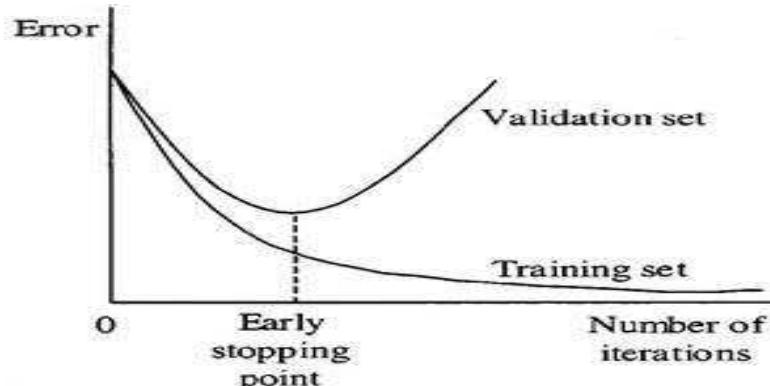


Better = AVOID overfitting by REGULARIZATION

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 41

Avoid overfitting by EARLY STOPPING

- For Neural Networks, a first method to avoid overfitting is to STOP LEARNING iterations as soon as the *validation_error* stops decreasing
- Generally, not a good idea to decide the number of iterations beforehand. Better to ALWAYS USE EARLY STOPPING



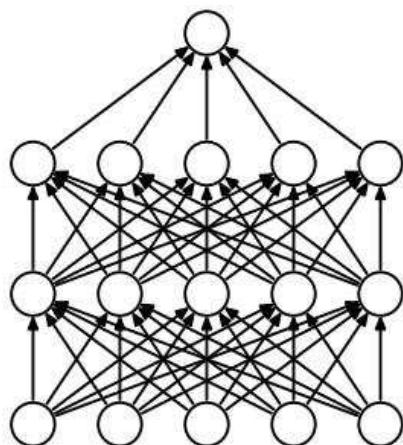
Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 42

Regularization = penalizing too complex models
Often done by adding a special term to cost function

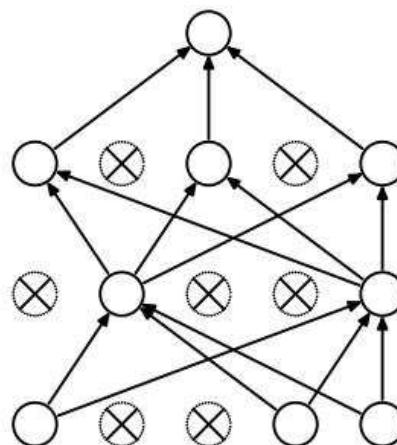
For neural network, the regularization term is just the L2- or L1- norm of the vector of all weights:

$$K = \sum_m (\text{loss}(Y_m, D_m)) + \beta \sum_{ij} |W_{ij}|^p \quad \text{with } p=2 \text{ (L2) or } p=1 \text{ (L1)}$$

→ name “Weight decay”



(a) Standard Neural Net



(b) After applying dropout.

At each training stage, individual nodes can be temporarily “dropped out” of the net with probability p (usually ~ 0.5), or re-installed with last values of weights

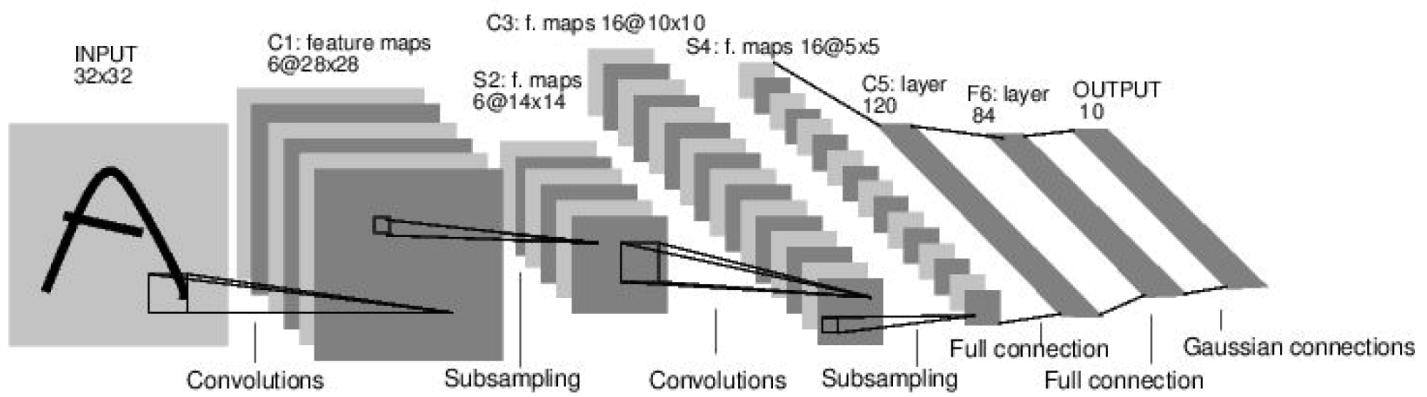
- **Introduction to Deep Learning**
- **Convolutional Neural Networks (CNN or ConvNets)**
 - Intro + Short reminder on Neural Nets
 - Convolution layers & Pooling layers + global architecture
 - Training algorithm + Dropout Regularization
- **Useful pre-trained convNets**
- **Coding frameworks**
- **Transfer Learning**
- **Object localization and Semantic segmentation**
- **Deep-Learning on 1D signal and 3D data**
- **Recent other image-based applications**

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 45

Examples of very successful ConvNets

- **LeNet:** 1st successful applications of ConvNets, by Yann LeCun in 1990's. Used to read zip codes, digits, etc.
- **AlexNet:** Beginning of ConvNet “buzz”: largely outperformed competitors in ImageNet_ILSVRC2012 challenge. Developped by Alex Krizhevsky et al., architecture similar to LeNet (but deeper+larger, and some chained ConvLayers before Pooling). 60 M parameters !
- **ZF Net:** ILSVRC 2013 winner. Developped by Zeiler&Fergus, by modif of AlexNet on some architecture hyperparameters.
- **GoogLeNet:** ILSVRC 2014 winner, developed by Google. Introduced an *Inception Module*, + AveragePooling instead of FullyConnected layer at output. Dramatic reduction of number of parameters (4M, compared to AlexNet with 60M).
- **VGGNet:** Runner-up in ILSVRC 2014. Very deep (16 CONV/FC layers) → 140M parameters !!
- **ResNet:** ILSVRC 2015, “Residual Network” introducing “skip” connections. Currently ~ SoA in convNet. Very long training but fast execution.

Input: 32x32 image



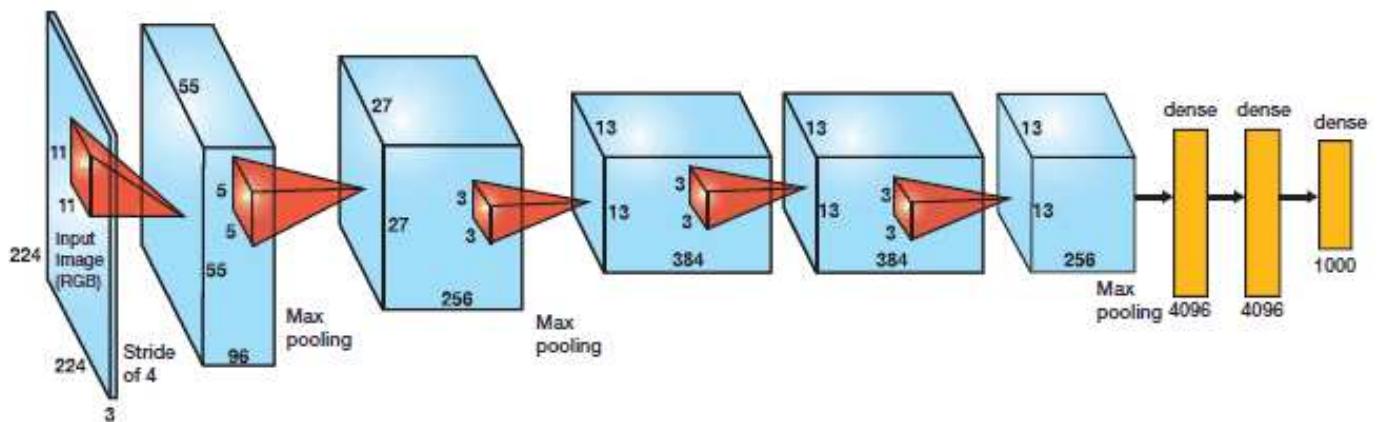
Conv filters were 5x5, applied at stride 1

Subsampling (Pooling) layers were 2x2 applied at stride 2

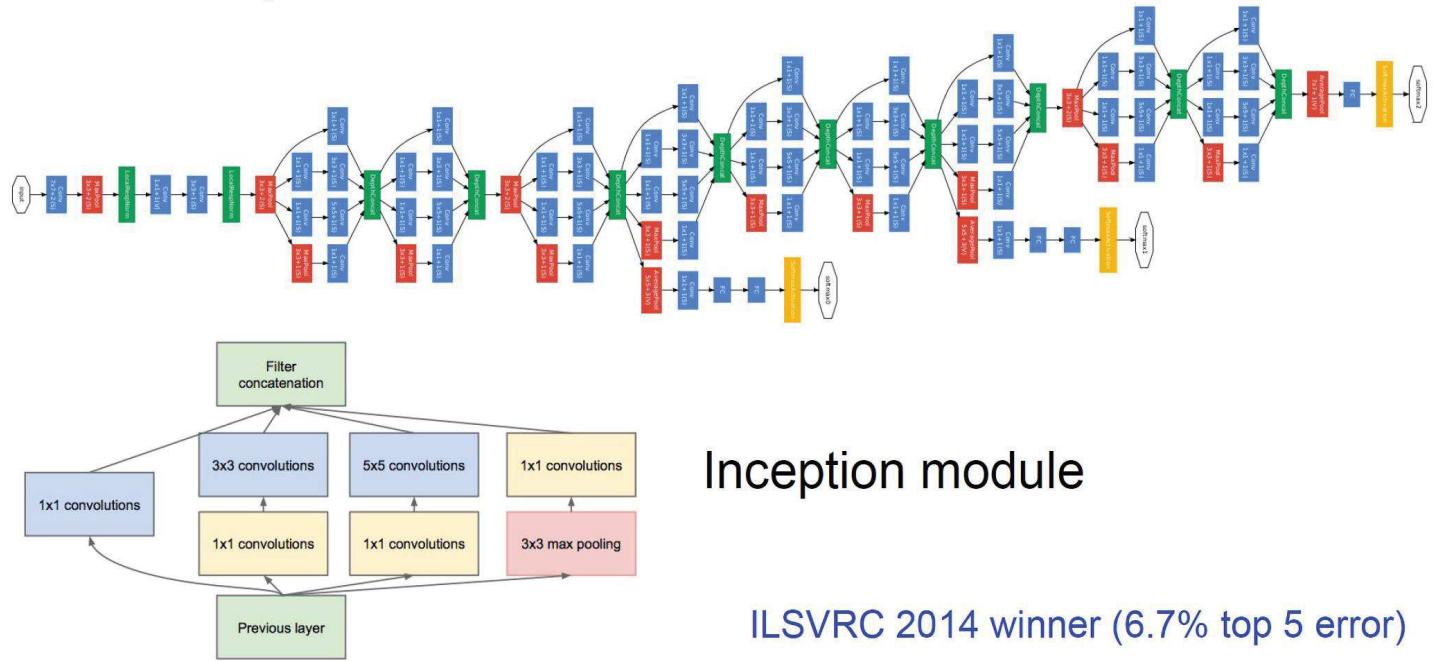
i.e. architecture is [CONV-POOL-CONV-POOL-CONV-FC]

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 47

Input: 224x224x3 image



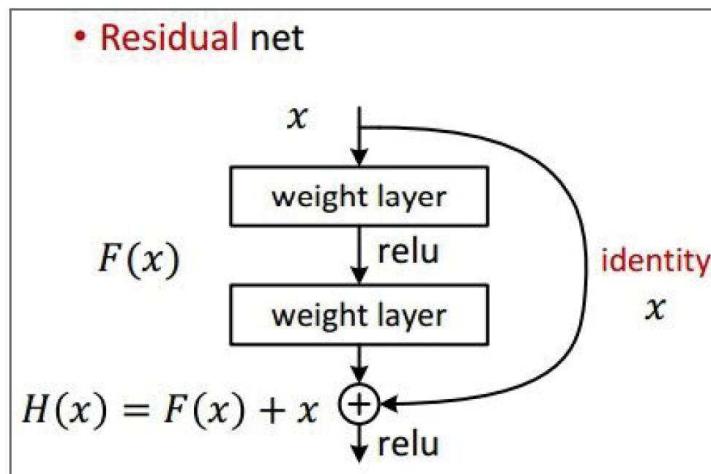
60 million parameters !...

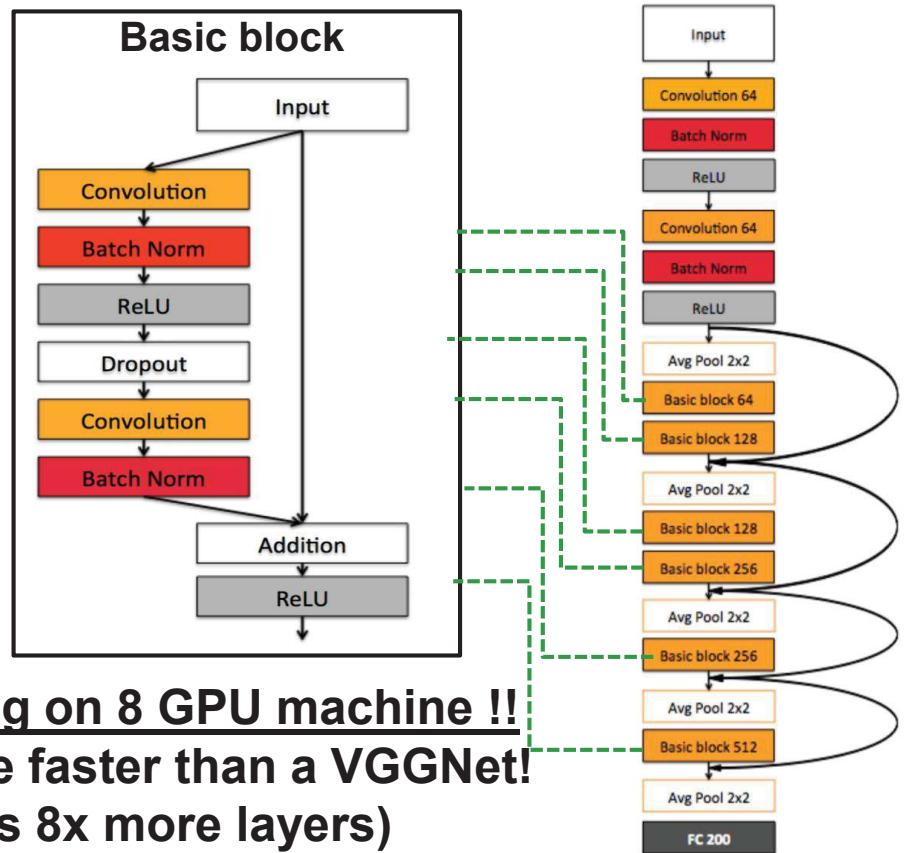


Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 50

ResNet (Residual Net), by Microsoft [He et al., 2015]

- ILSVRC 2015 large winner in 5 main tracks (3.6% top 5 error)
- 152 layers!!!
- But novelty = "skip" connections

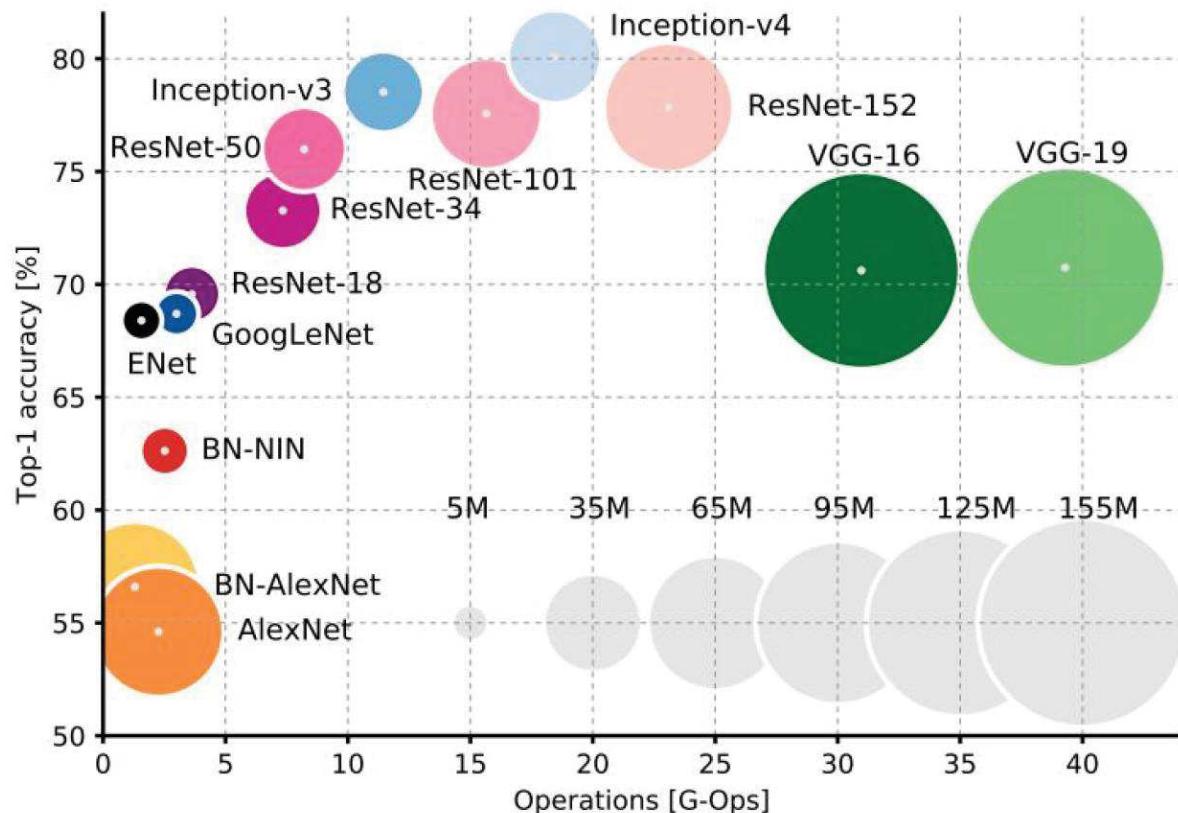


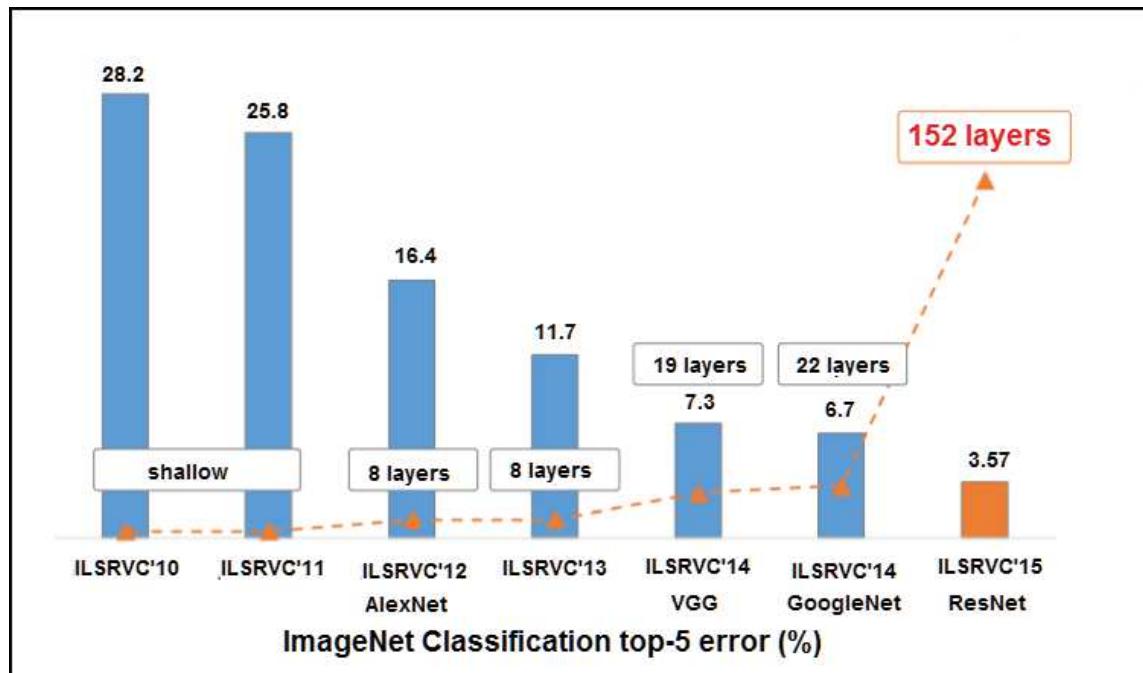


- 2-3 weeks of training on 8 GPU machine !!**
- However, at runtime faster than a VGGNet!
(even though it has 8x more layers)

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 52

Performance comparison of usual convNet architectures





**But most important is the choice of
ARCHITECTURAL STRUCTURE**

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 54

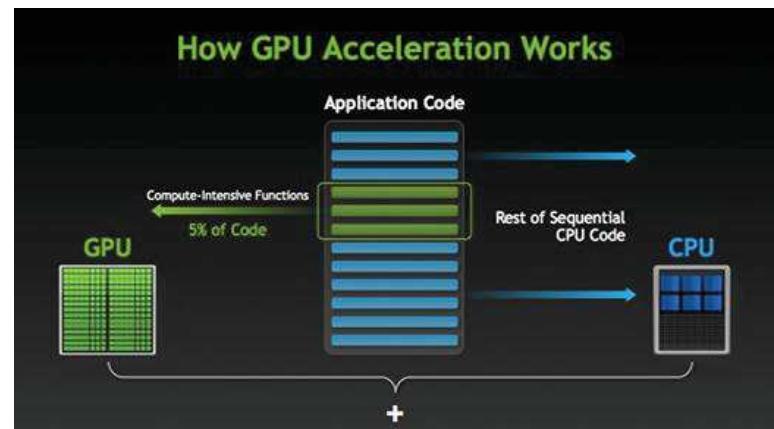
Outline

- Introduction to Deep Learning
- Convolutional Neural Networks (CNN or ConvNets)
 - Intro + Short reminder on Neural Nets
 - Convolution layers & Pooling layers + global architecture
 - Training algorithm + Dropout Regularization
- Useful pre-trained convNets
- Coding frameworks
- Transfer Learning
- Object localization and Semantic segmentation
- Deep-Learning on 1D signal and 3D data
- Recent other image-based applications

Good convNets are very big (millions of parameters!)

Training generally performed on BIG datasets

→ Training time more manageable using GPU acceleration for ultra-parallel processing



Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 56

Programming environments for Deep-Learning

- **TensorFlow** <https://www.tensorflow.org>
- **KERAS** <https://keras.io>
Python front-end APIs mapped either
on Tensor-Flow or Theano back-end
- **PyTorch** <https://pytorch.org/>
- **Caffe** <http://caffe.berkeleyvision.org/>
C++ library, hooks from Python → notebooks
- **Theano** <http://wwwdeeplearning.net/software/theano/>
- **Lasagne** <http://lasagne.readthedocs.io>
lightweight library to build+train neural nets in Theano

All of them handle transparent use of GPU,
and most of them are used in Python code/notebook

```

model = Sequential()

# 1 set of (Convolution+Pooling) layers, with Dropout
model.add(Convolution2D(conv_depth_1, kernel_size, kernel_size,
                       border_mode='valid', input_shape=(depth, height, width)))
model.add( MaxPooling2D(pool_size=(pooling_size, pooling_size)) )
model.add(Activation('relu'))
model.add(Dropout(drop_prob))

# Now flatten to 1D, and apply 1 Fully_Connected layer
model.add(Flatten())
model.add(Dense(hidden_size1, init='lecun_uniform'))
model.add(Activation('sigmoid'))

# Finally add a Softmax output layer, with 1 neuron per class
model.add(Dense(num_classes, init='lecun_uniform'))
model.add(Activation('softmax'))

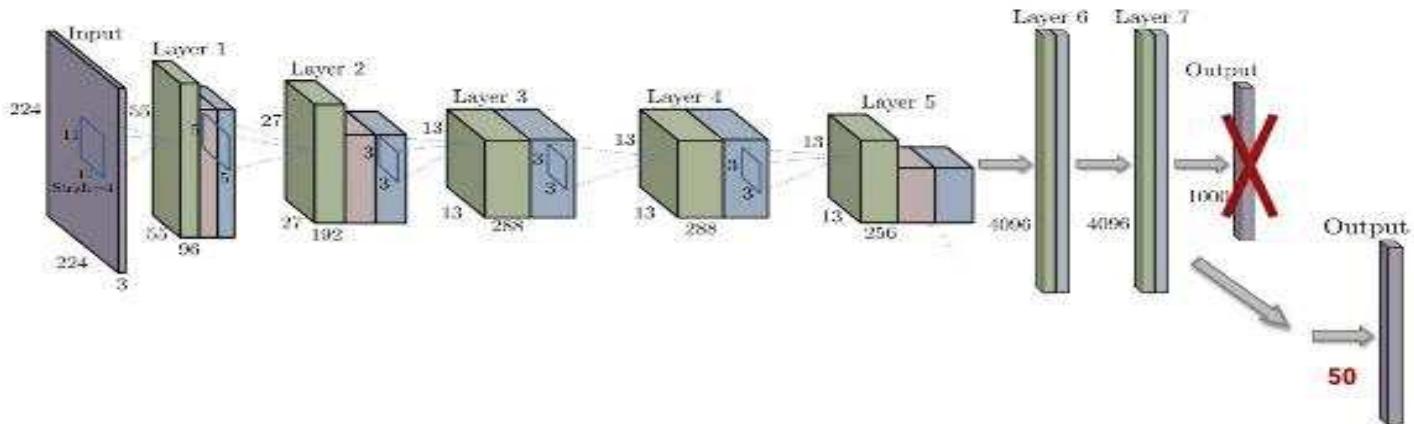
# Training "session"
sgd = SGD(lr=learning_rate, momentum=0.8) # Optimizer
model.compile(loss='categorical_crossentropy', optimizer=sgd)
model.fit(X_train, Y_train, batch_size=32, nb_epoch=2, verbose=1,
           validation_split=valid_proportion)

# Evaluate the trained model on the test set
model.evaluate(X_test, Y_test, verbose=1)

```

Outline

- **Introduction to Deep Learning**
- **Convolutional Neural Networks (CNN or ConvNets)**
 - Intro + Short reminder on Neural Nets
 - Convolution layers & Pooling layers + global architecture
 - Training algorithm + Dropout Regularization
- **Useful pre-trained convNets**
- **Coding frameworks**
- **Transfer Learning**
- **Object localization and Semantic segmentation**
- **Deep-Learning on 1D signal and 3D data**
- **Recent other image-based applications**

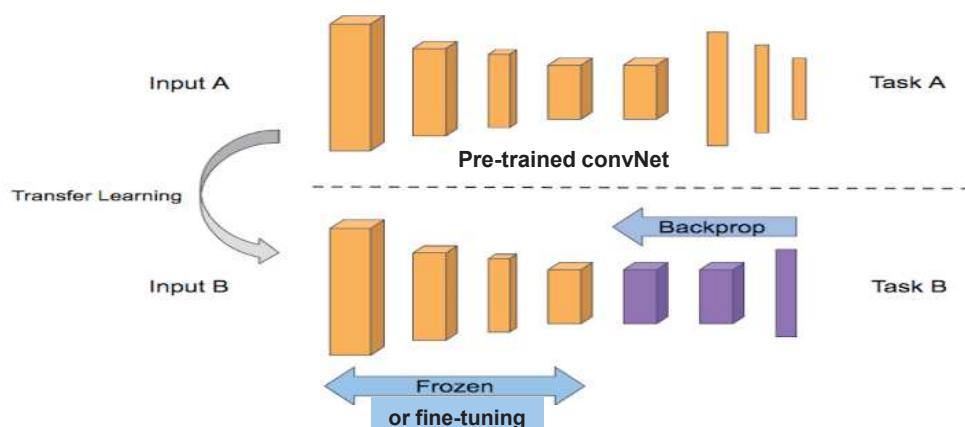


By removing last layer(s) (those for classification) of a convNet trained on ImageNet, one obtains a transformation of any input image into a semi-abstract representation, which can be used for learning SOMETHING ELSE (« transfer learning »):

- either by just using learnt representation as features
- or by creating new convNet output and perform learning of new output layers + fine-tuning of re-used layers

Transfer learning and fine-tuning

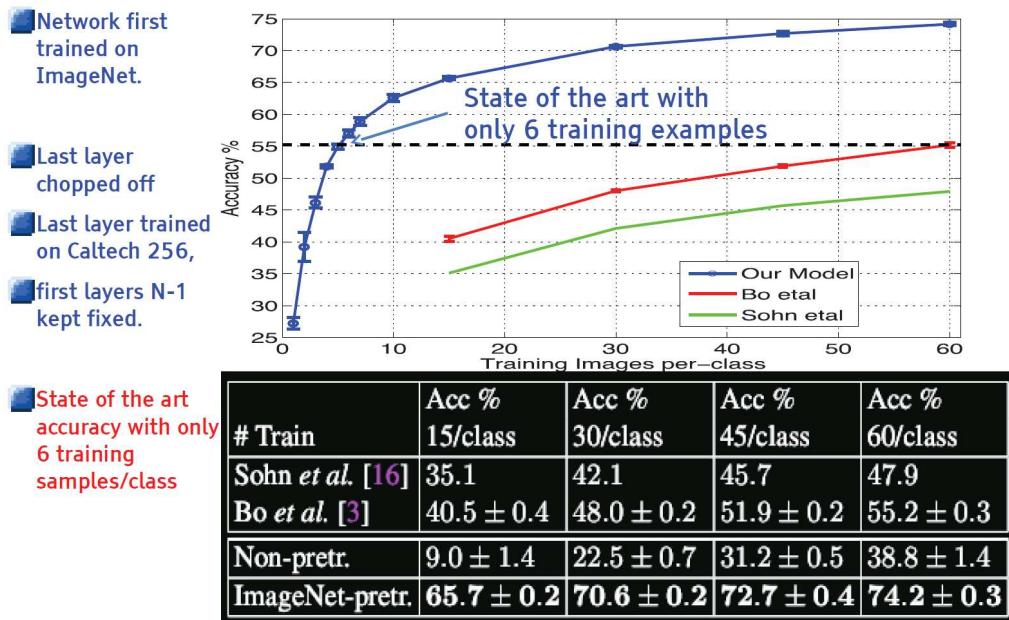
- SoA convNets trained on ImageNet are image CLASSIFIERS for one object per image
- Many object categories can be irrelevant (e.g. boat in a office)
- For each application, models are usually obtained from state-of-the-art ConvNets pre-trained on ImageNet (winners of yearly challenge, eg: AlexNet, VGG, Inception, ResNet, etc...)



→ Adaptation is performed by Transfer Learning, ie modification+training of last layers and/or fine-tuning of pre-trained weights of lower layers

Transfer Learning with few training examples

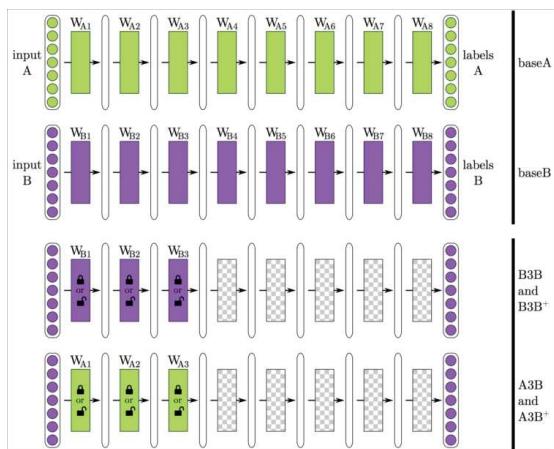
- Using a CNN pre-trained on a large dataset, possible to adapt it to another task, using only a SMALL training set!



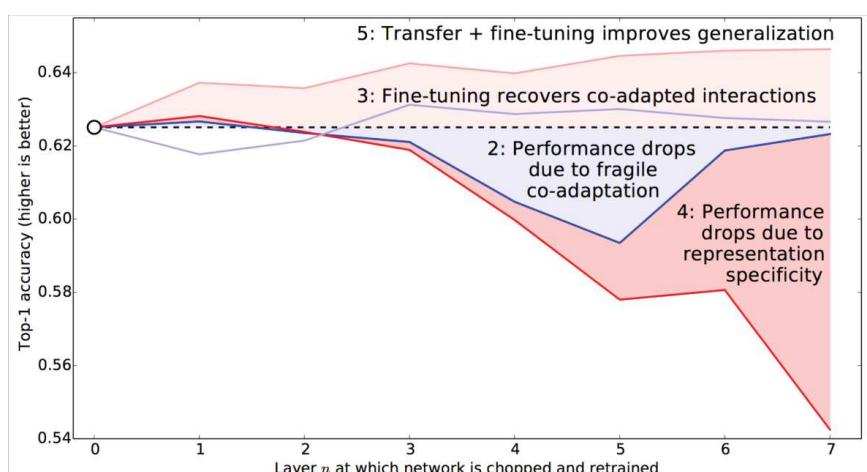
3: [Bo, Ren, Fox. CVPR, 2013] 16: [Sohn, Jung, Lee, Hero ICCV 2011]

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 62

Transfer-Learning even improves performances!



[Yosinski, Clune, Bengio, Lipson,
"How transferable are features in
deep neural networks?", ICML'2014]



Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 63

- Learning on simulated synthetic images
+ fine-tuning on real-world images
 - Recognition/classification for OTHER categories or classes
 - Training an objects detector (or a semantic segmenter)
-
- Precise localization (position+bearing) = PoseNet
 - Human posture estimation = openPose
 - End-to-end driving (imitation Learning)
 - 3D informations (depth map) from monovision!

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 64

Transfer Learning code example in Keras

```
from keras.applications.inception_v3 import InceptionV3
from keras.preprocessing import image
from keras.models import Model
from keras.layers import Dense, GlobalAveragePooling2D
from keras import backend as K
# create the base pre-trained model
base_model = InceptionV3(weights='imagenet',
                           include_top=False)

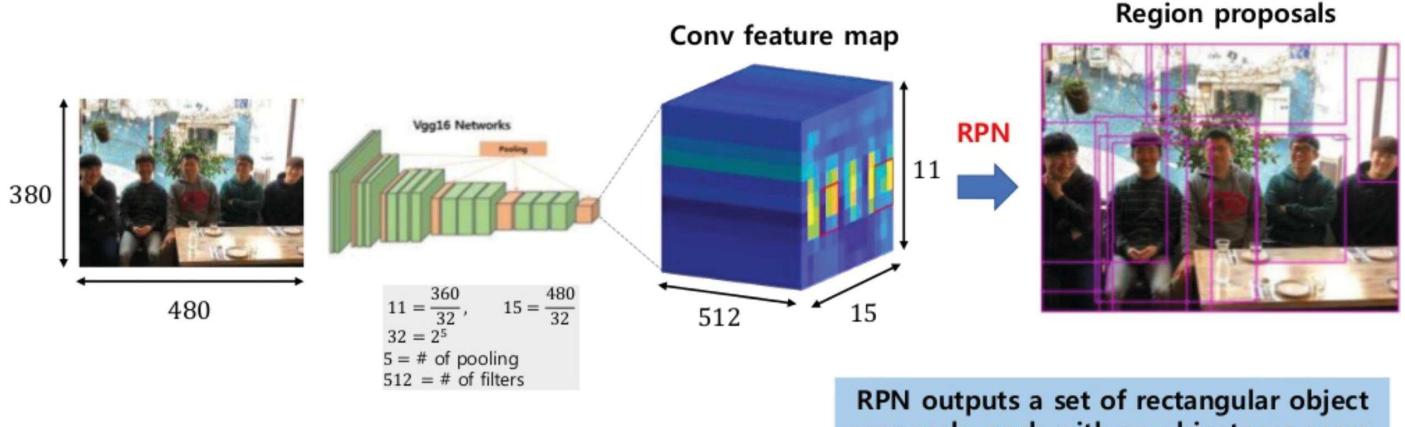
# add a global spatial average pooling layer
x = base_model.output
x = GlobalAveragePooling2D()(x)
# let's add a fully-connected layer
x = Dense(1024, activation='relu')(x)
# and a logistic layer -- let's say we have 200 classes
predictions = Dense(200, activation='softmax')(x)
# this is the model we will train
model = Model(input=base_model.input, output=predictions)
# first: train only the top layers (which were randomly initialized)
# i.e. freeze all convolutional InceptionV3 layers
for layer in base_model.layers:
    layer.trainable = False
# compile the model (should be done *after* setting layers to non-trainable)
model.compile(optimizer='rmsprop', loss='categorical_crossentropy')
# train the model on the new data for a few epochs
model.fit_generator(...)
```

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 65

- Introduction to Deep Learning
- Convolutional Neural Networks (CNN or ConvNets)
 - Intro + Short reminder on Neural Nets
 - Convolution layers & Pooling layers + global architecture
 - Training algorithm + Dropout Regularization
- Useful pre-trained convNets
- Coding frameworks
- Transfer Learning
- Object localization and Semantic segmentation
- Deep-Learning on 1D signal and 3D data
- Recent other image-based applications

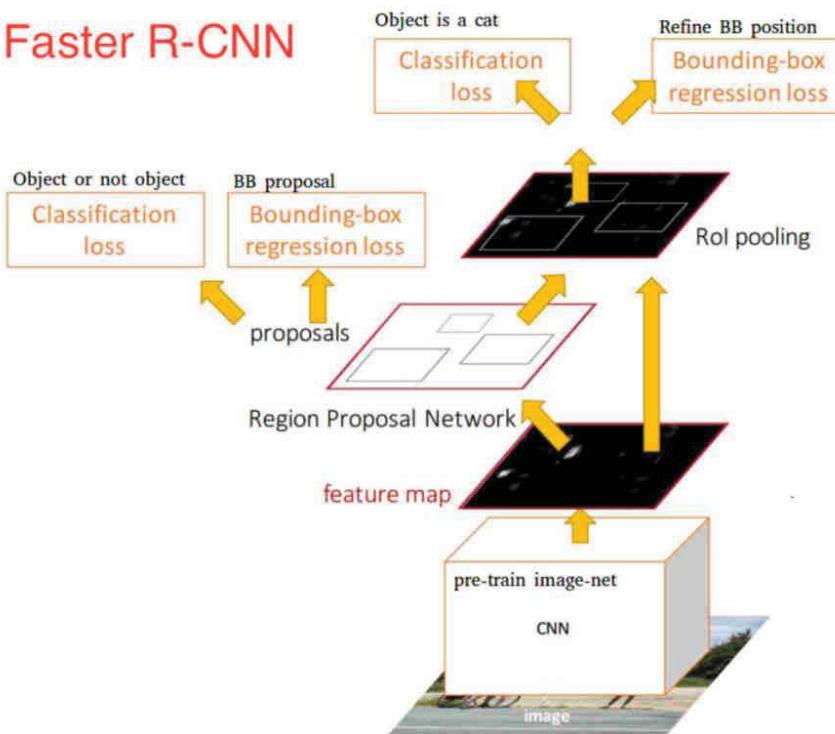
Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 66

Deep-Learning for visual object DETECTION



The high-level representation computed by last convolution layer can be analyzed for detection and localization (bounding-boxes) of all objects of interesting categories

Faster R-CNN



**Region Proposal Network (RPN) on top of standard convNet.
End-to-end training with combination of 4 losses**

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 68

Example of visual DETECTION & categorization with Faster_R-CNN



**ConvNets are currently state-of-the-art
ALSO for visual objects detection**

Solve detection as a regression problem ("single-shot" detection)

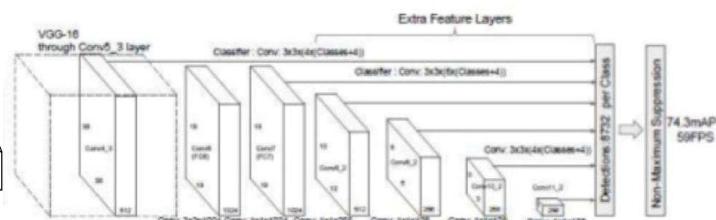
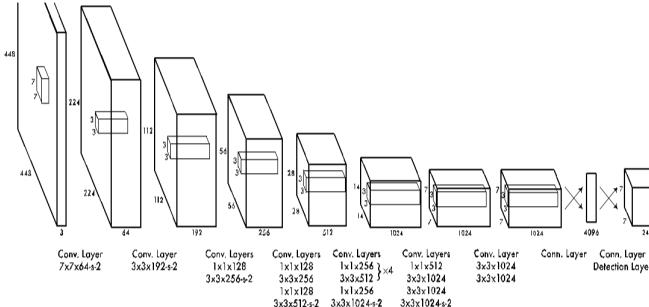
YOLO

and

SSD

YOU ONLY LOOK ONCE(YOLO)

SINGLE SHOT MULTIBOX DETECTOR(SSD)

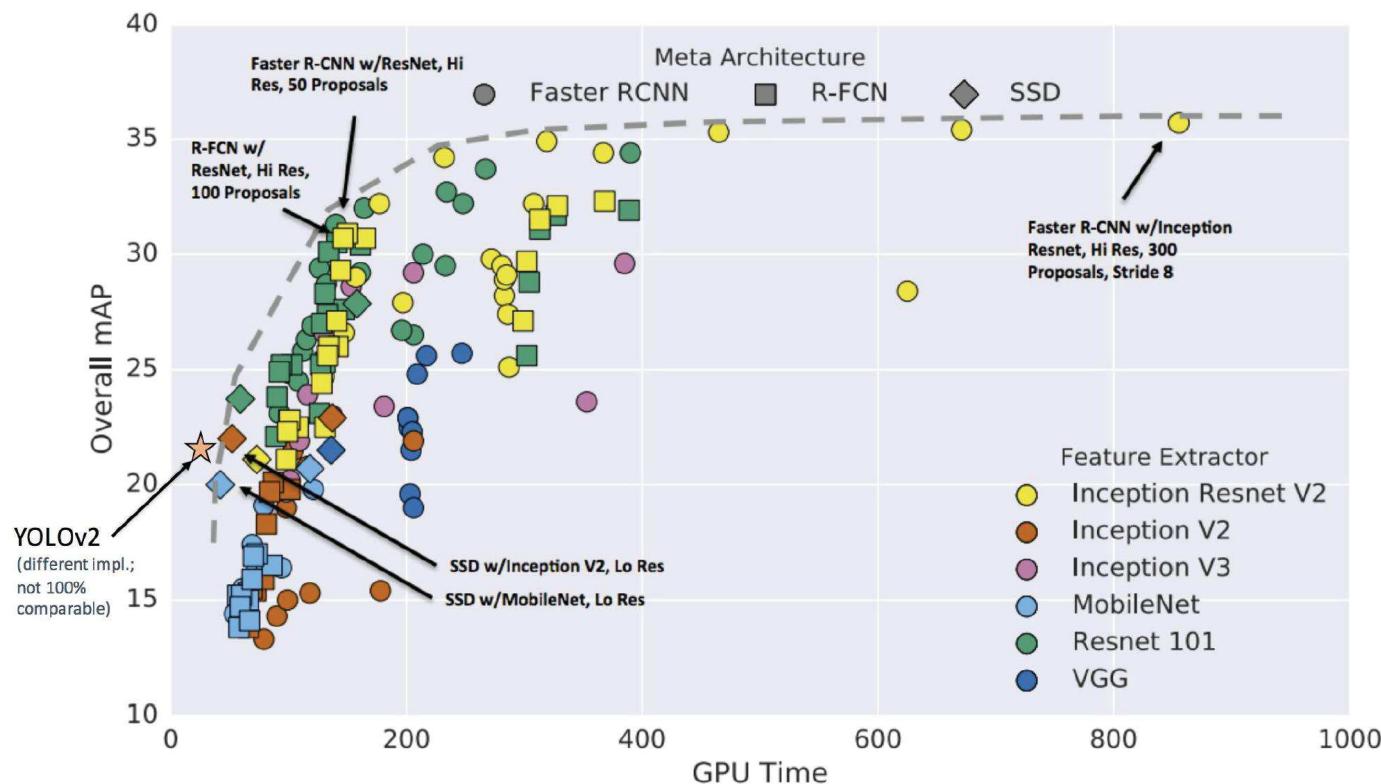


Images from: <https://www.slideshare.net/TaegyunJeon1/pr12-you-only-look-once-yolo-unified-realtime-object-detection>

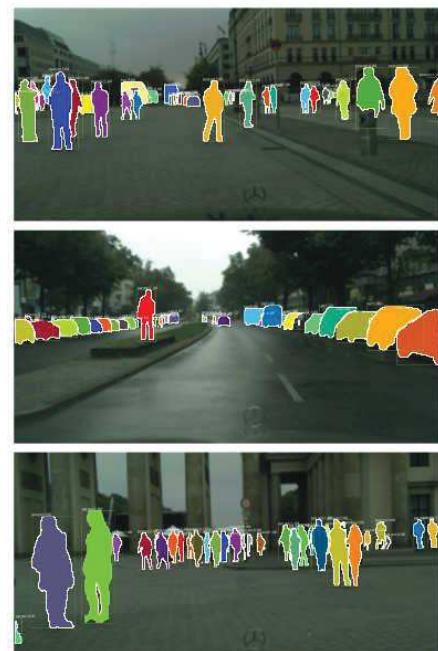
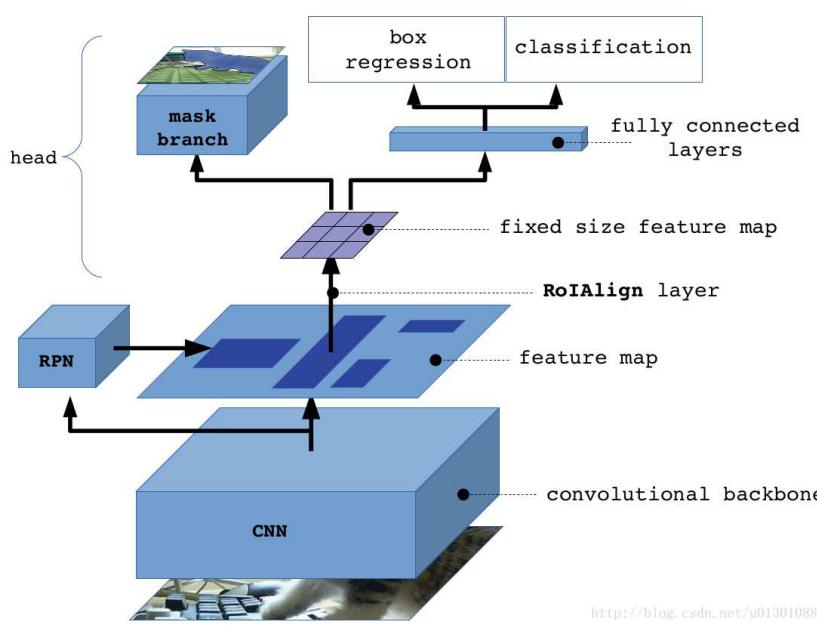
Both are faster, but less accurate, than Faster_R-CNN

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 70

Recent comparison of object detection convNets



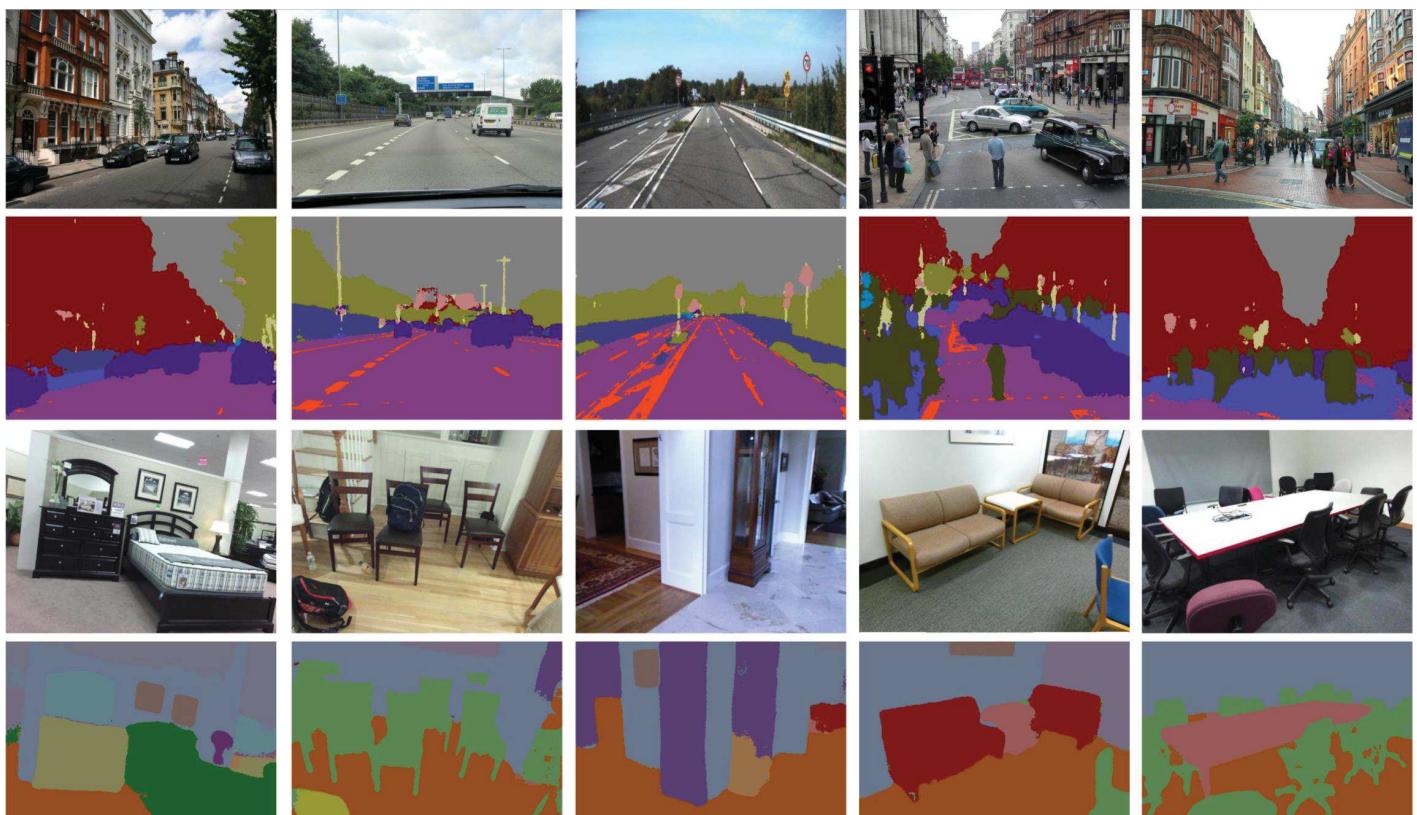
Slide from Ross Girshick's [CVPR 2017 Tutorial](#), Original Figure from Huang et al



Mask R-CNN architecture (left) extracts detailed contours and shape of objects instead of just bounding-boxes

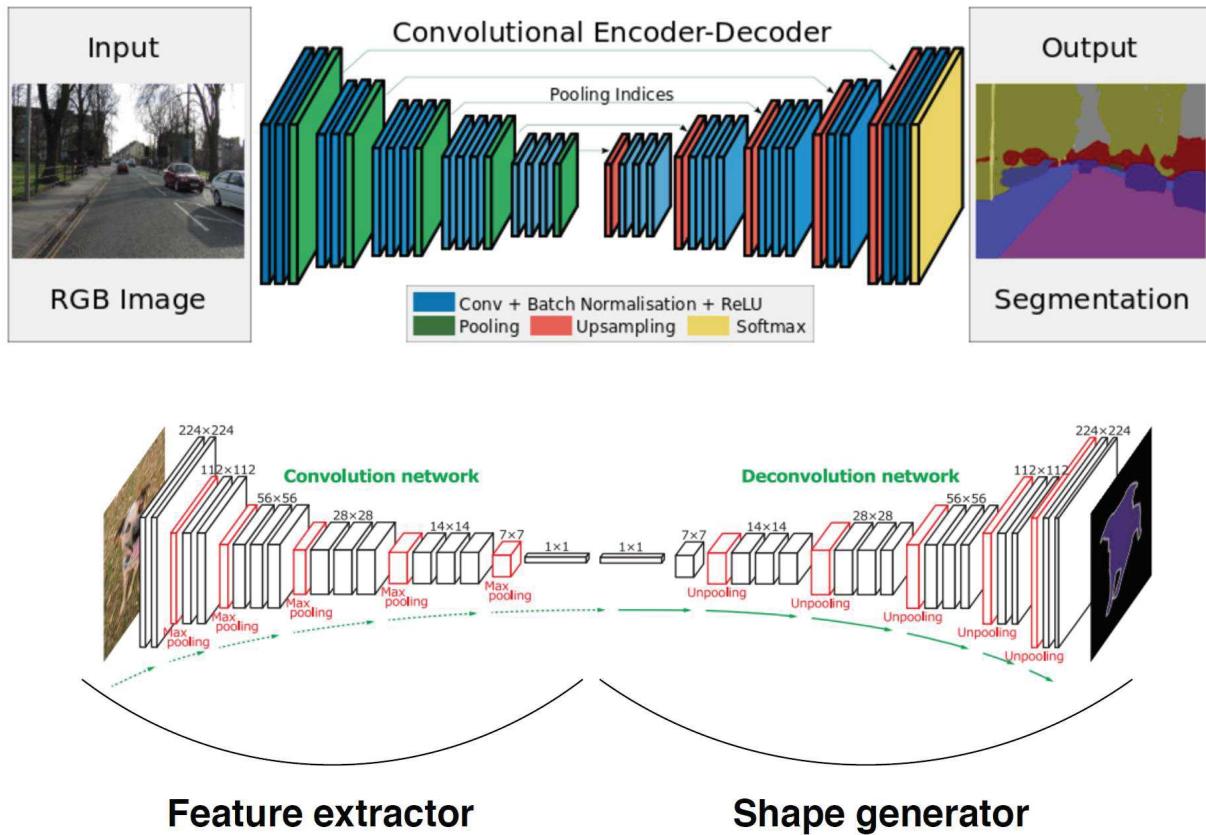
Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 72

Semantic segmentation



Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 73

Convolutional Encoder-Decoder



Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 74

Many competitors for DL of semantic segmentation

Many competitors for semantic segmentation by deep-learning:

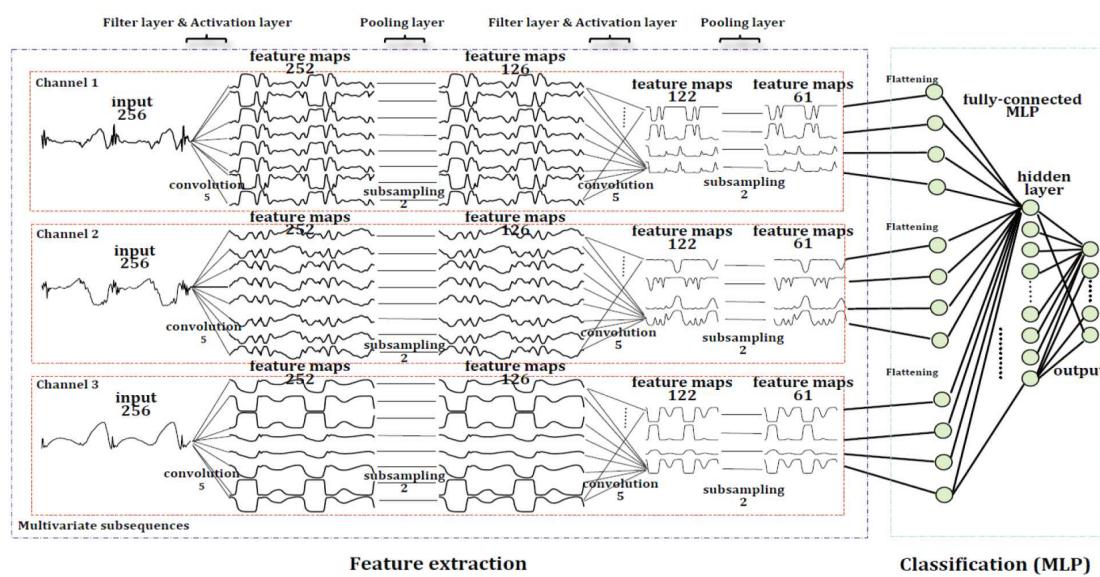
- SegNet (2015)
- U-Net (2015)
- RefineNet (2016)
- ICnet (2017)
- DeepLab
- ...

VERY HOT TOPIC !!!

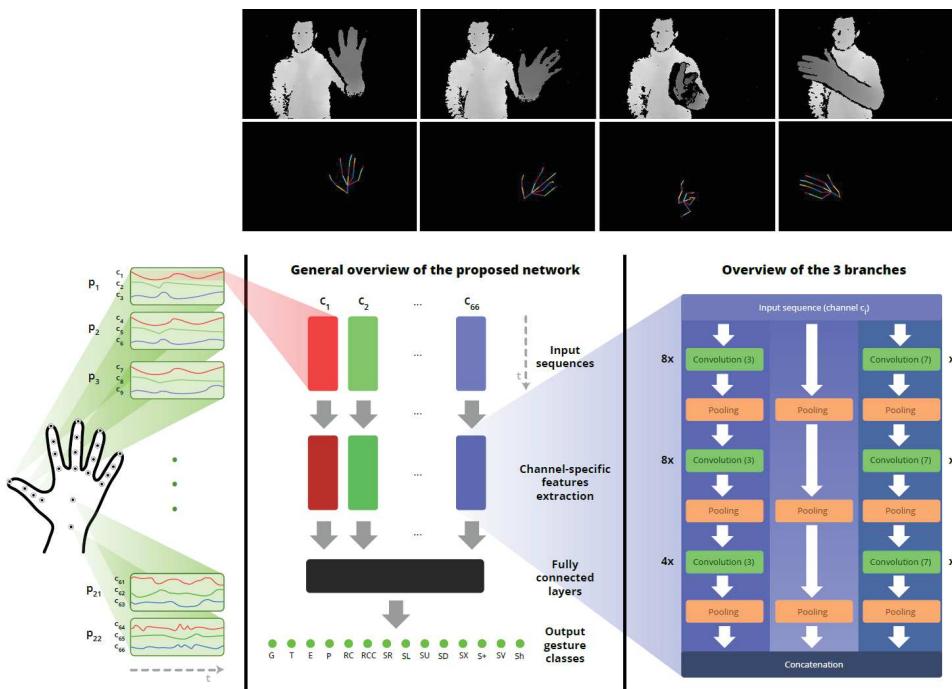
- Introduction to Deep Learning
- Convolutional Neural Networks (CNN or ConvNets)
 - Intro + Short reminder on Neural Nets
 - Convolution layers & Pooling layers + global architecture
 - Training algorithm + Dropout Regularization
- Useful pre-trained convNets
- Coding frameworks
- Transfer Learning
- Object localization and Semantic segmentation
- Deep-Learning on 1D signal and 3D data
- Recent other image-based applications

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 76

Deep-TEMPORAL Convolution for multivariate *time-series*



**MC-DCNN model
(separate 1D temporal convolution of each time-serie)**



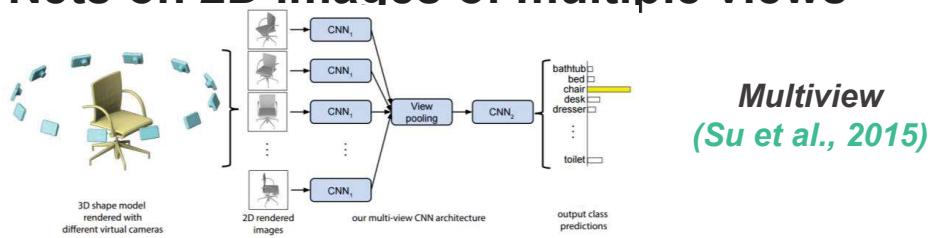
Work in progress at center for Robotics of MINES ParisTech
(PhD thesis of Guillaume Devineau)

Potential applicability to other kinds of time-series!

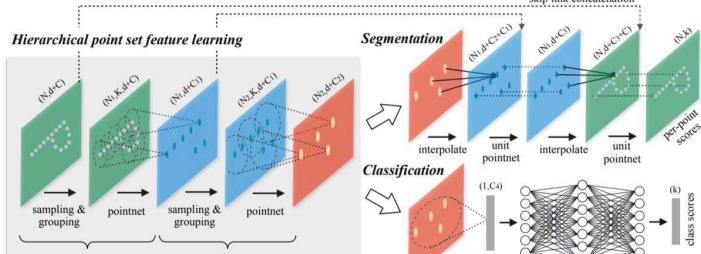
Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 78

Possible to use:

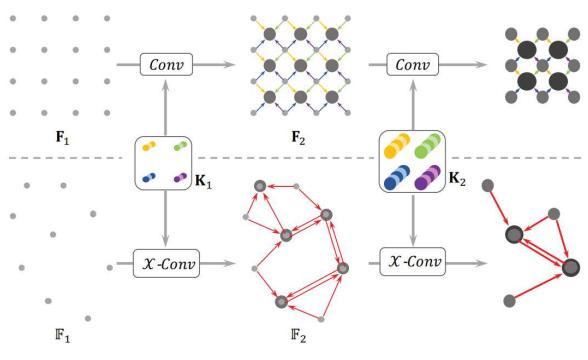
- ConvNets on 2D images of multiple views



- ConvNet on 2D DEPTH image(s)
 - Convolutions of 3D points

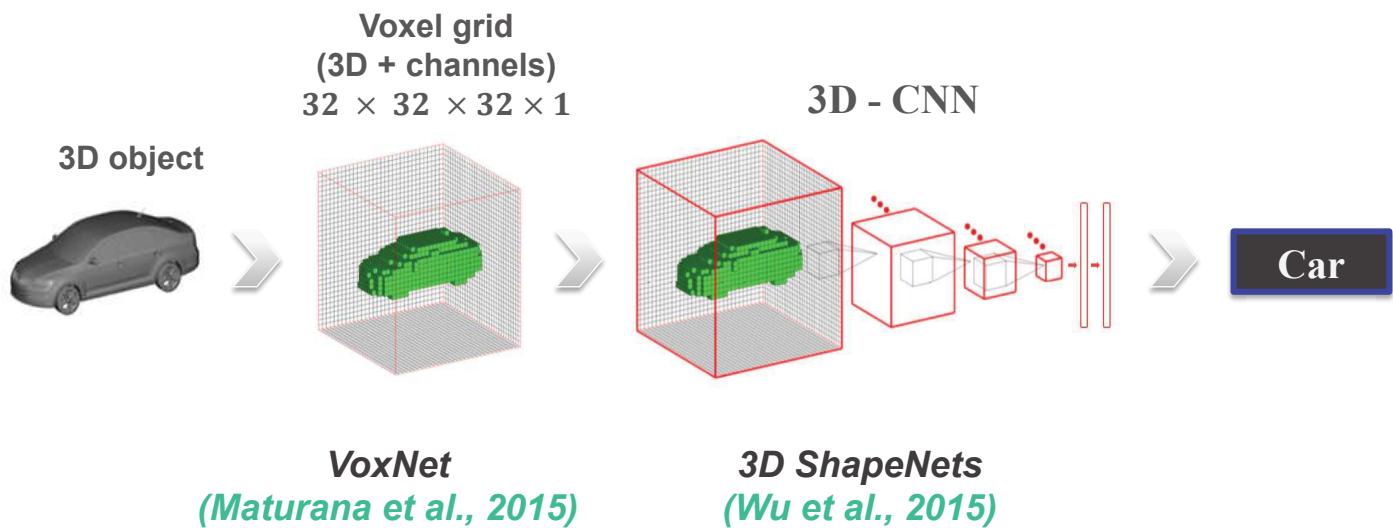


PointNet++ (Qi et al., 2017)



PointCNN (Li et al., 2018)

- 3D convolutions on voxels (see next slide)



Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 80

Outline

- **Introduction to Deep Learning**
- **Convolutional Neural Networks (CNN or ConvNets)**
 - Intro + Short reminder on Neural Nets
 - Convolution layers & Pooling layers + global architecture
 - Training algorithm + Dropout Regularization
- **Useful pre-trained convNets**
- **Coding frameworks**
- **Transfer Learning**
- **Object localization and Semantic segmentation**
- **Deep-Learning on 1D signal and 3D data**
- **Recent other image-based applications**

- Image classification
- Visual object detection and categorization
- Semantic segmentation of images
- ...

AND ALSO:

- Image-based localization
- Estimation of Human pose
- Inference of 3D (depth) from monocular vision
- Learning image-based behaviors
 - End-to-end driving from front camera
 - Learning robot behavior from demonstration/imitation

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 82

PoseNet: 6-DoF camera-pose regression with Deep-Learning



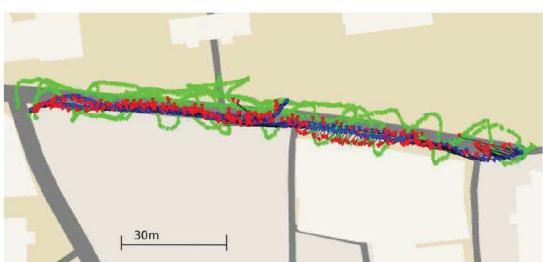
Input RGB Image



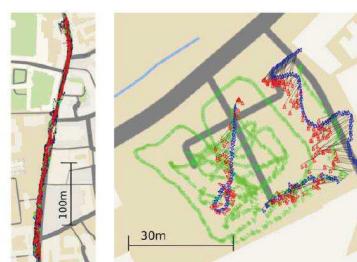
Convolutional
Neural Network
(GoogLeNet)



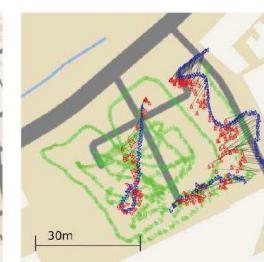
6-DOF
Camera Pose



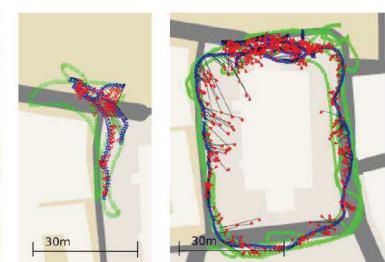
King's College



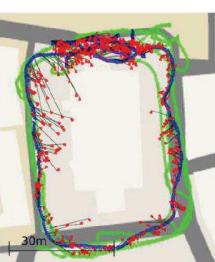
Street



Old Hospital



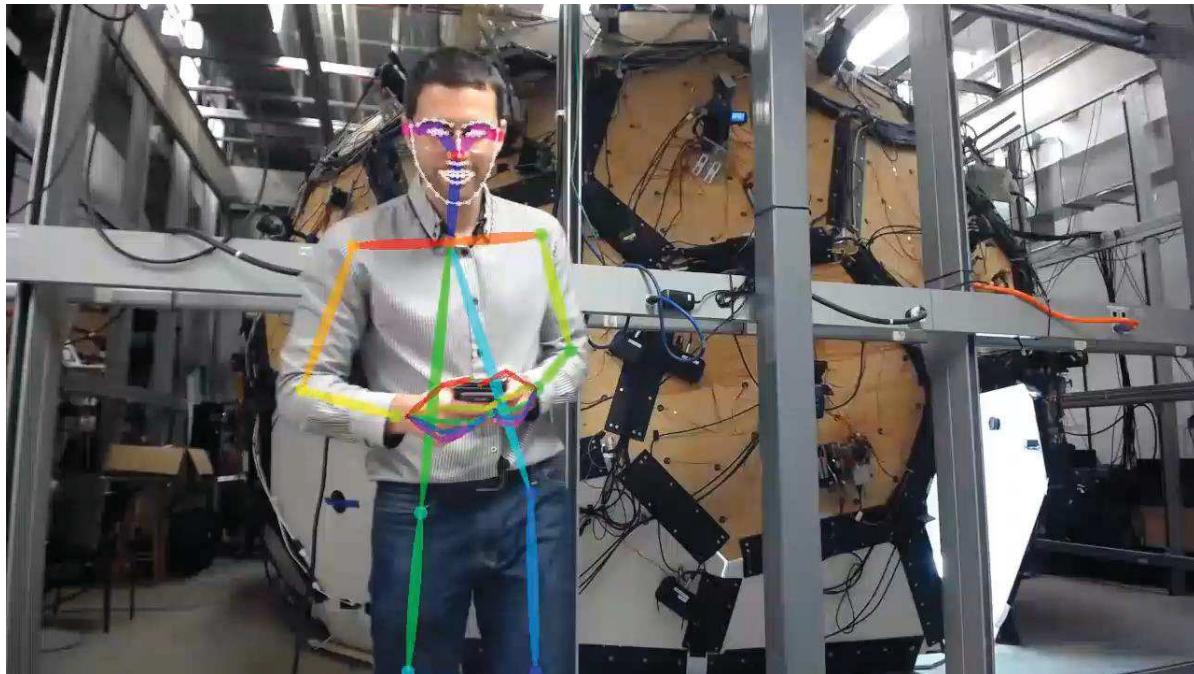
Shop Façade



St Mary's Church

Figure 4: Map of dataset showing training frames (green), testing frames (blue) and their predicted camera pose (red). The testing sequences are distinct trajectories from the training sequences and each scene covers a very large spatial extent.

[A. Kendall, M. Grimes & R. Cipolla, "PoseNet: A Convolutional Network for Real-Time 6-DOF Camera Relocalization", ICCV'2015, pp. 2938-2946]

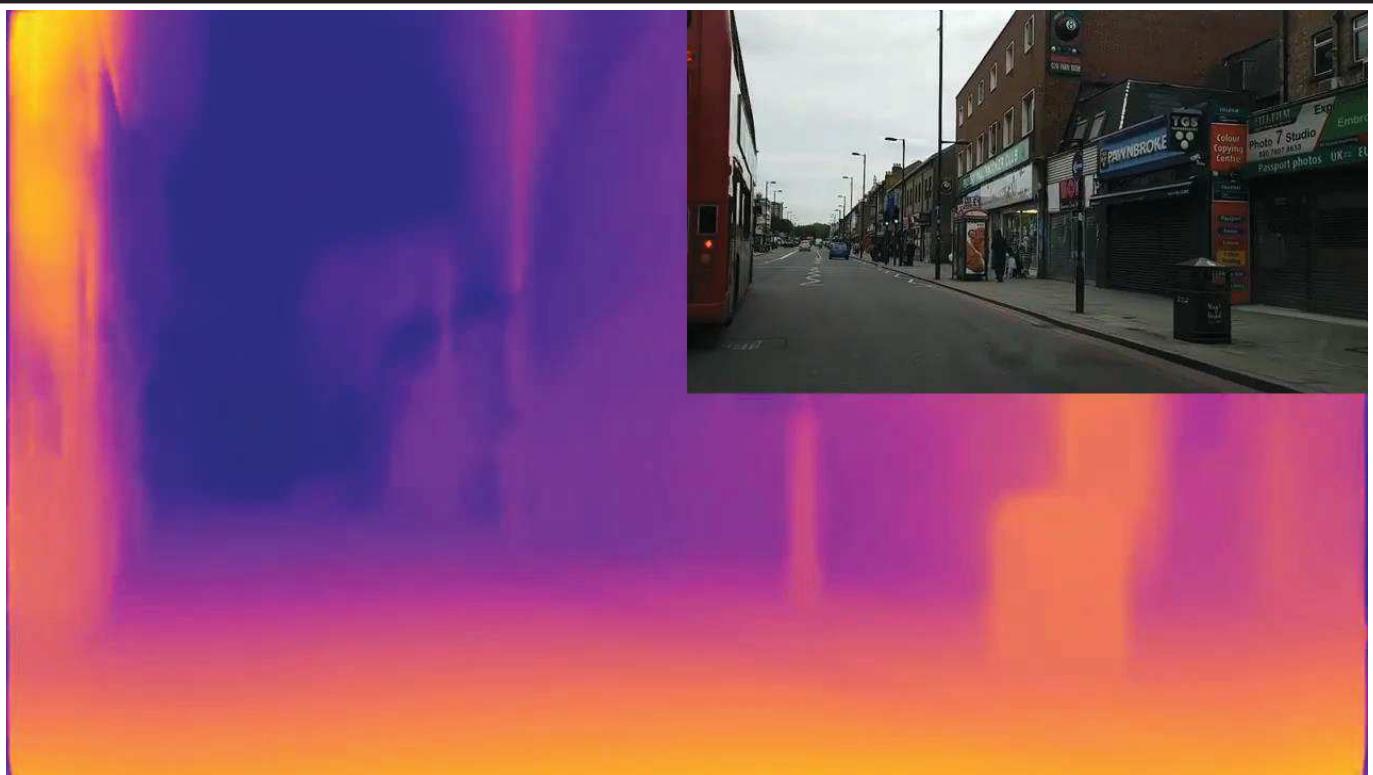


Real-time estimation of Human poses on RGB video

OpenPose [*Realtime Multi-Person 2D Pose Estimation using Part Affinity Field*,
Cao et al., CVPR'2017 [CMU]]

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 84

Inference of 3D (depth) from monocular vision

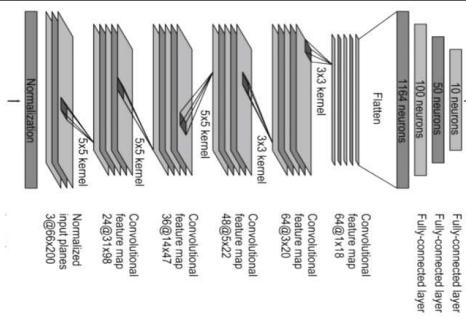


Unsupervised monocular depth estimation with left-right consistency
C Godard, O Mac Aodha, GJ Brostow - CVPR'2017 [UCL]

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 85

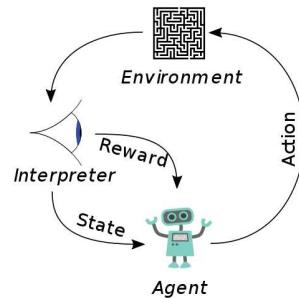


ConvNet input:
Cylindrical projection of
fisheye camera



ConvNet output:
steering angle

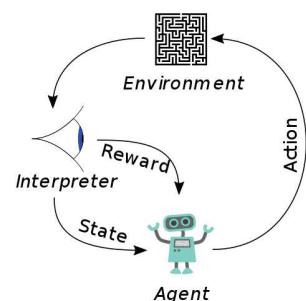
Imitation Learning from Human driving on real data

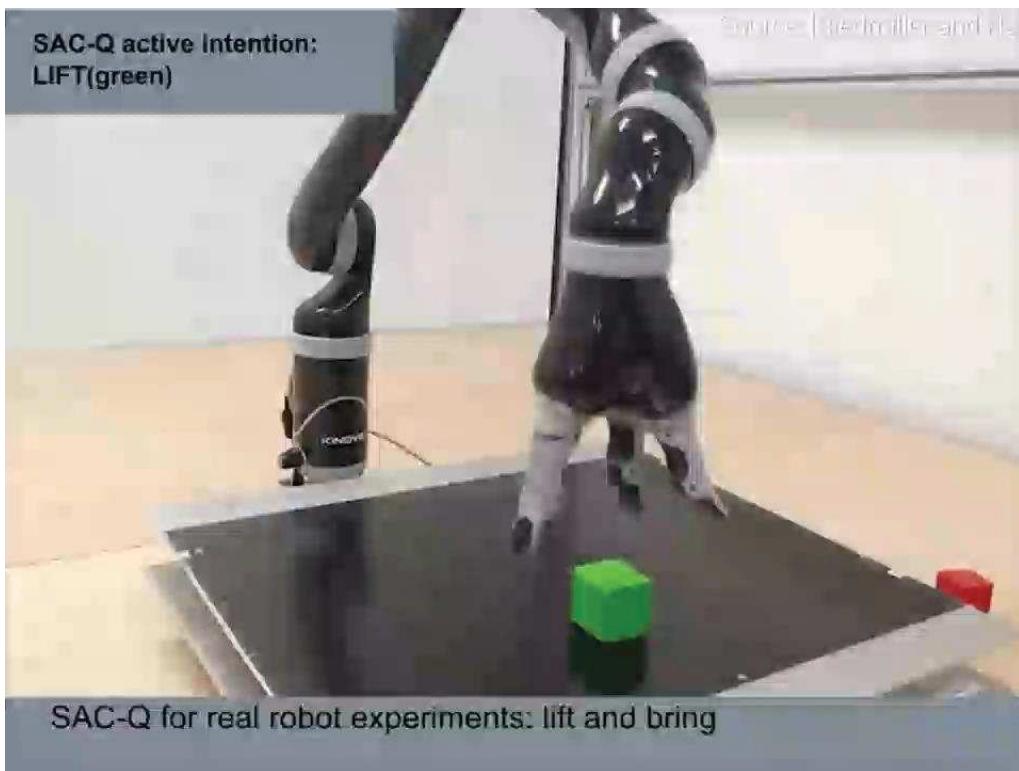


End-to-end driving via Deep Reinforcement Learning
[thèse CIFRE Valeo/MINES-ParisTech en cours]

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 86

Robot task learning using Reinforcement Learning





Work by Google DeepMind

[*Learning by Playing Solving Sparse Reward Tasks from Scratch, Riedmiller et al. (ICML'2018)*]

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 88

Summary on ConvNets & Deep-Learning

- Proven advantage of learning features empirically from data
- Large ConvNets require huge amounts of labelled examples data for training
- Current research/progresses = finding efficient global architecture of ConvNets
- Enormous potential of TRANSFER-LEARNING on small datasets for restricted/specialized problems
- ConvNets also for multivariate time-series (1D temporal convolutions) and for 3D data (3D conv on voxels, etc...)
- ConvNets can potentially infer from image ANYTHING for which information is in the image (3D, movement, planning, ...)

Next frontiers:

- Theoretical aspects
- Robustness issues (cf. adversarial examples)



- UNsupervised deep-learning on unlabelled data
- Deep Reinforcement Learning (DRL)
- Deep Recurrent Neural Networks (LSTM, GRU, etc...) for sequence processing (NLP!) or modeling behavior & dynamics

Deep-Learning: general principles + convNets, Pr. Fabien MOUTARDE, Center for Robotics, MINES ParisTech, PSL, Nov.2019 90

Any QUESTIONS ?